

Addressing Multiphysics and Hardware Complexity in a HPC Environment

Tony Saad and James C. Sutherland
University of Utah

Workshop on Software Development Environments for High-Performance Computing
October 18, 2015 • San Francisco, CA

DOE Awards
DE-NA0002375
DE-NA-000740



National Science Foundation
WHERE DISCOVERIES BEGIN
NSF PetaApps award 0904631



Matt Might



James Sutherland



Hao Hou



Michael Ballantyne



Christopher Earl



Tony Saad



Abhishek Bagusetty

DOE Awards
DE-NA0002375
DE-NA-000740



National Science Foundation
WHERE DISCOVERIES BEGIN
NSF PetaApps award 0904631

Three Challenges

Three Challenges

1. Hardware Complexity

Three Challenges

1. Hardware Complexity



Three Challenges

2. Programming Complexity

Three Challenges

2. Programming Complexity

```
for (int i = 0; i<nx; ++i) {
    for (int j=0; j<ny; ++j) {
        for(int k=0;k<nz; ++k){
            rhs = (phi[i+1, j, k] - 2* phi[i,j,k] + phi[i-1,j,k])/dx2
                  + (phi[i, j+1, k] - 2* phi[i,j,k] + phi[i,j-1,k])/dy2
                  + (phi[i, j, k+1] - 2* phi[i,j,k] + phi[i,j,k-1])/dz2;
        }
    }
}
```

Three Challenges

3. Multiphysics Complexity

Three Challenges

3. Multiphysics Complexity



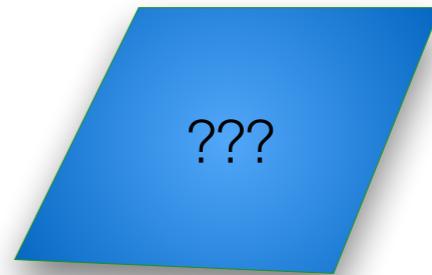
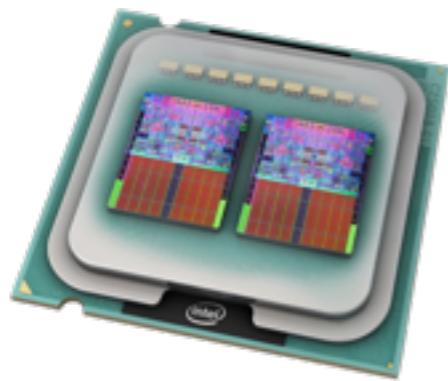
Three Challenges

1. Hardware Complexity
2. Programming Complexity
3. Multiphysics Complexity

Three Solutions

**write one code
execute
on arbitrary architectures**

**write one code
execute
on arbitrary architectures**



write expressive syntax

$$\nabla \cdot (k \nabla \phi)$$

write expressive syntax

$$\nabla \cdot (k \nabla \phi)$$

div(k * grad(phi))

write expressive syntax

$$\nabla \cdot (k \nabla \phi)$$

`div(k * grad(phi))`

that is as fast as nested loops

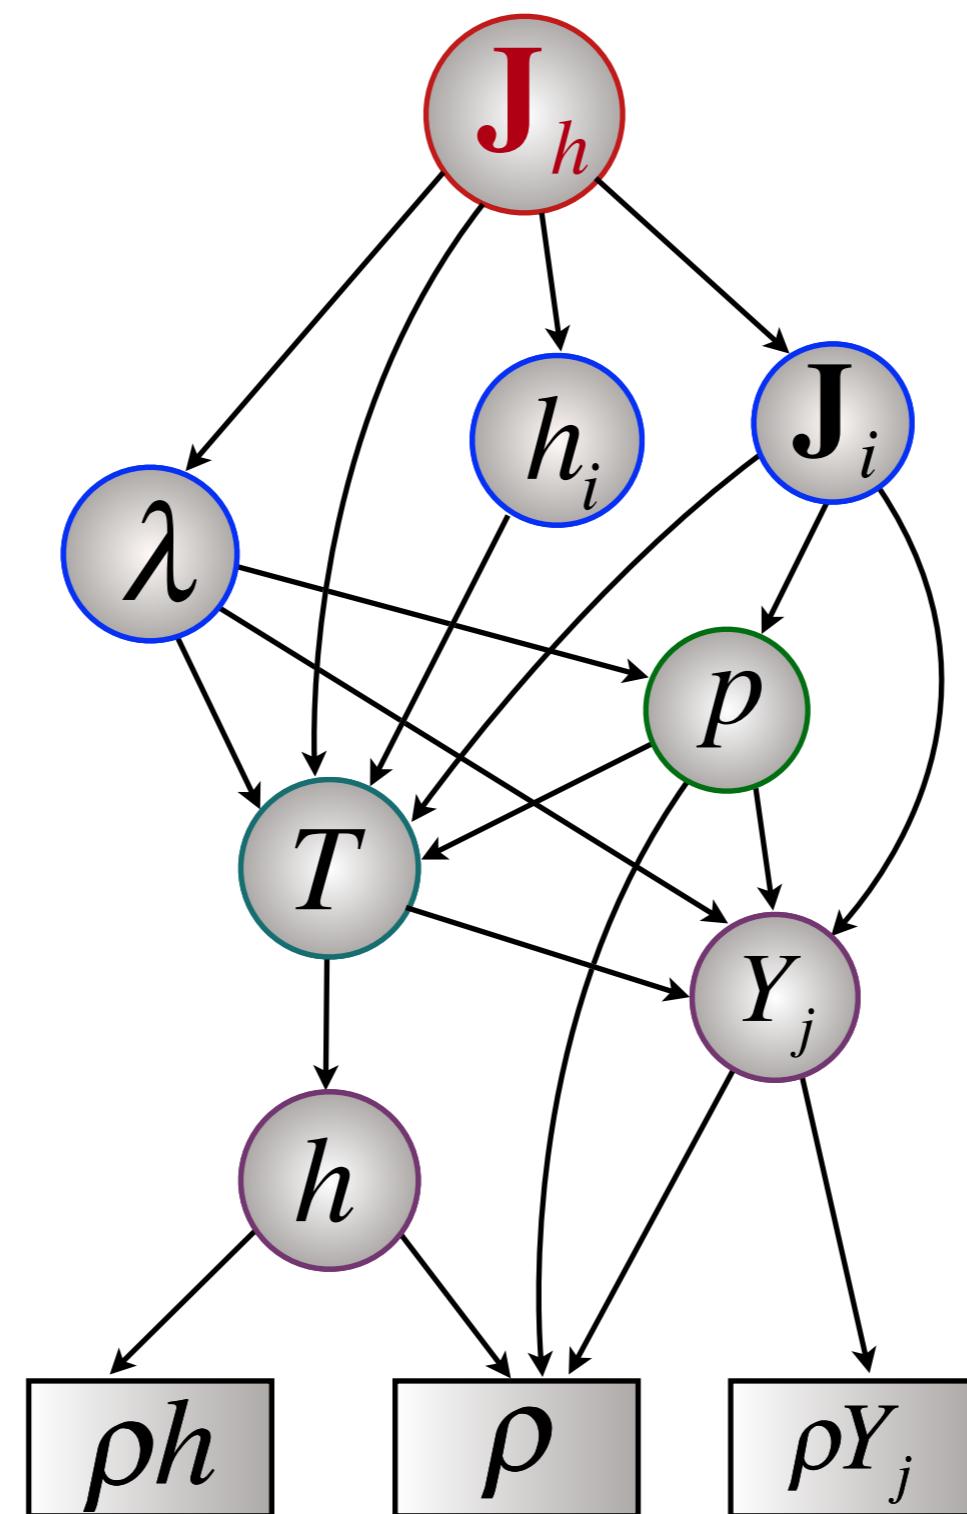
that is **as fast as** nested loops

div(k * grad(phi))

```
for (int i = 0; i<nx; ++i) {
    for (int j=0; j<ny; ++j) {
        for(int k=0;k<nz; ++k){
            rhs = (phi[i+1, j, k] - 2* phi[i,j,k] + phi[i-1,j,k])/dx2
                  + (phi[i, j+1, k] - 2* phi[i,j,k] + phi[i,j-1,k])/dy2
                  + (phi[i, j, k+1] - 2* phi[i,j,k] + phi[i,j,k-1])/dz2;
        }
    }
}
```

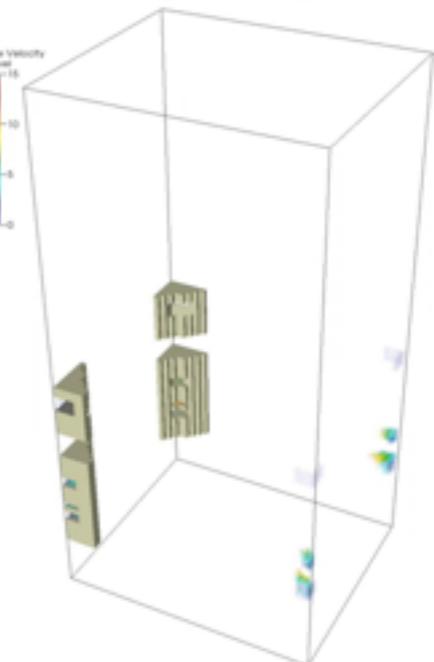
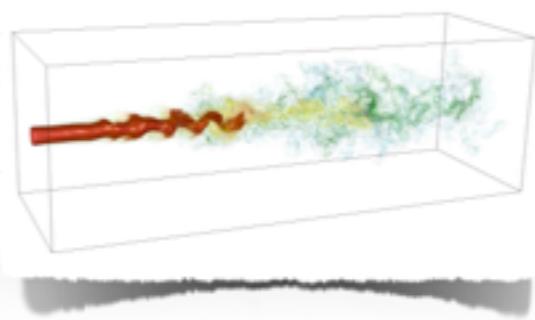
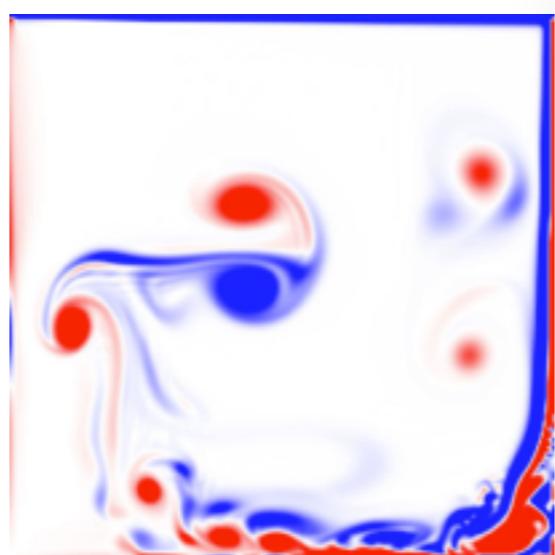
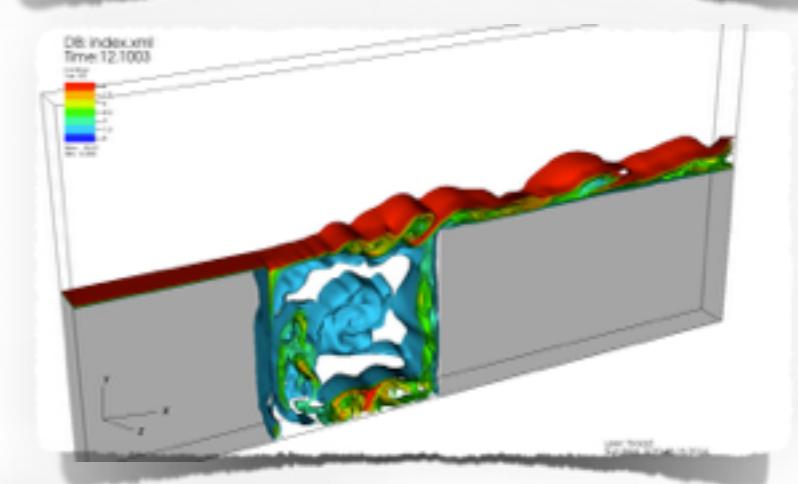
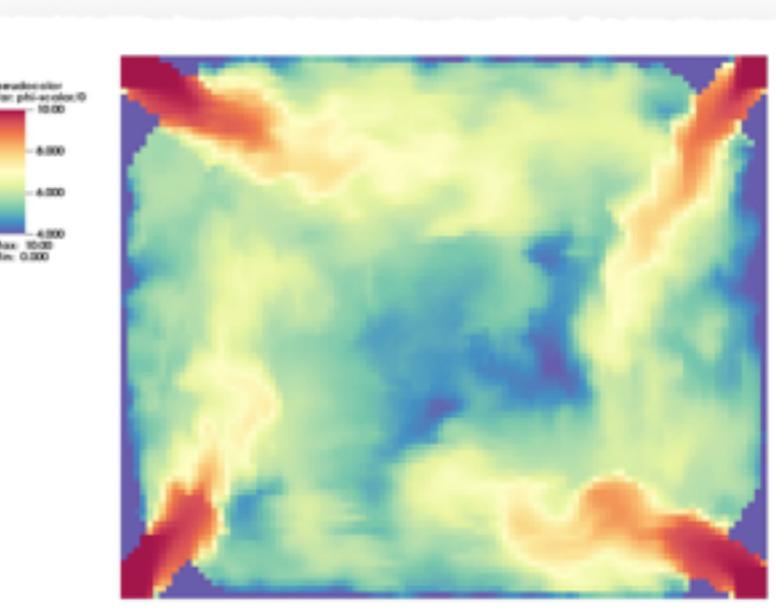
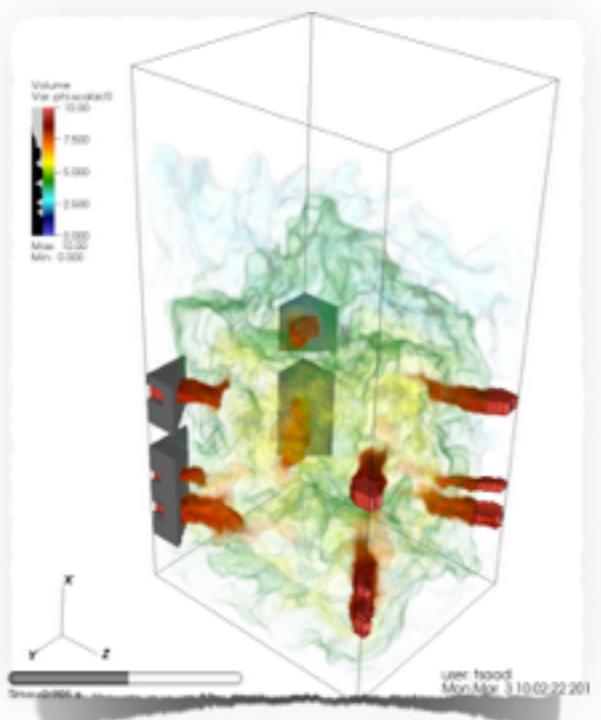
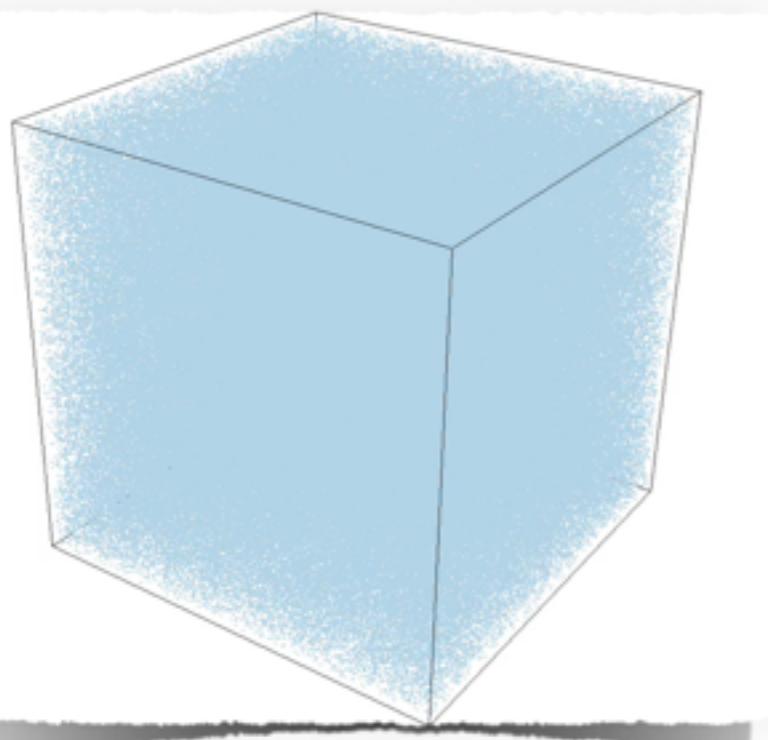
Tame multi-physics complexity

Tame multi-physics complexity



Wasatch
(et al.)

Wasatch is
a multi-architecture
CFD
and multi-physics code
that enables fast deployment
of physics models
and PDEs



Can't do it alone...

Life in an HPC Ecosystem

Wasatch

ExprLib

Nebo/
SpatialOps

Uintah

Uintah

- Task based parallel infrastructure
- Provides parallel IO
- Grid specification
- Domain decomposition

Wasatch

ExprLib

Nebo/
SpatialOps

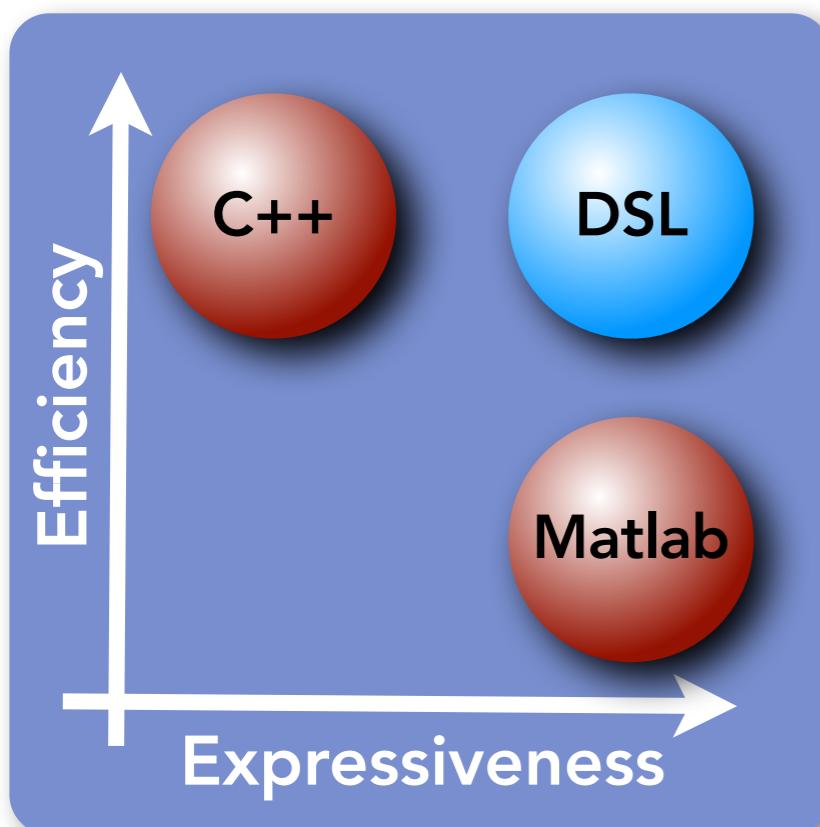
Uintah

Nebo DSL

Field & stencil operations:
(and much more)

```
rhs <= -divOpX( xConvFlux + xDiffFlux )  
      -divOpY( yConvFlux + yDiffFlux )  
      -divOpZ( zConvFlux + zDiffFlux );
```

Can “chain” stencil operations.



- 70+ natively supported discrete operators, easily define new ones
- Masks - allow operations on field subsets
- CPU (serial, multicore), GPU, Xeon Phi backends
- conditional operations (vectorized ‘if’)

ExprLib

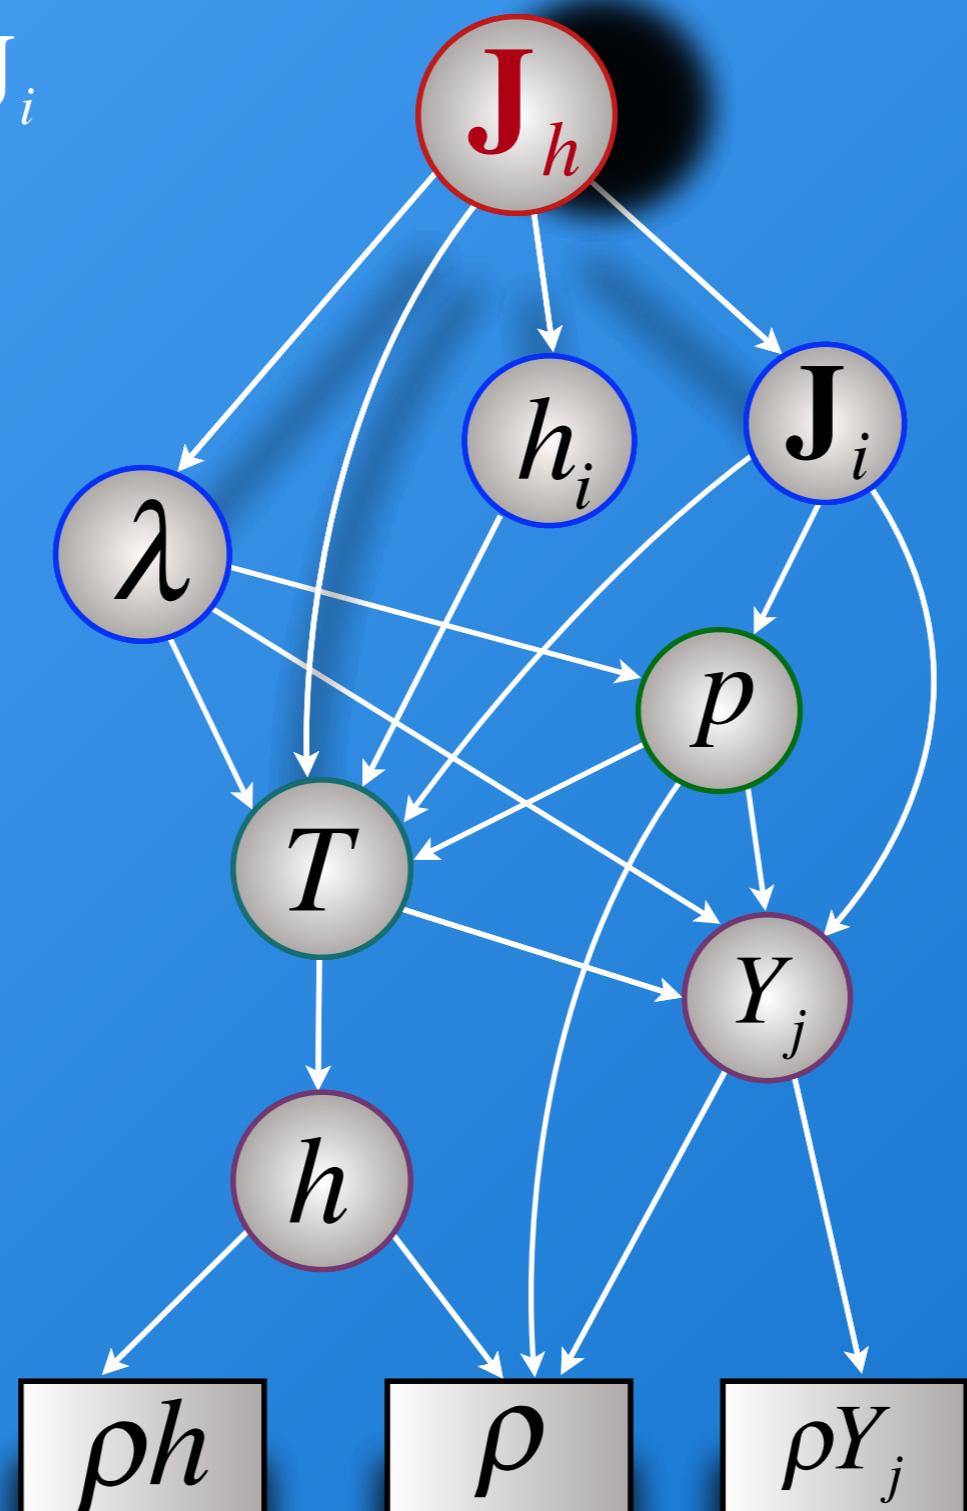
$$\mathbf{J}_h = -\lambda \nabla T + \sum_{i=1}^{n_s} h_i \mathbf{J}_i$$

MODEL A

$$\lambda = \lambda_0 = \text{const}$$

$$\mathbf{J}_i = -\sum_{j=1}^{n_s} D \nabla Y_j$$

$$h_i = h_i(T)$$



MODEL B

$$\lambda = \lambda(T, p, Y_i)$$

$$\mathbf{J}_i = -\sum_{j=1}^{n_s} D_{ij}(T, p, Y_k) \nabla Y_j$$

$$- D_i^T(T, p, Y_k) \nabla T$$

$$h_i = h_i(T)$$

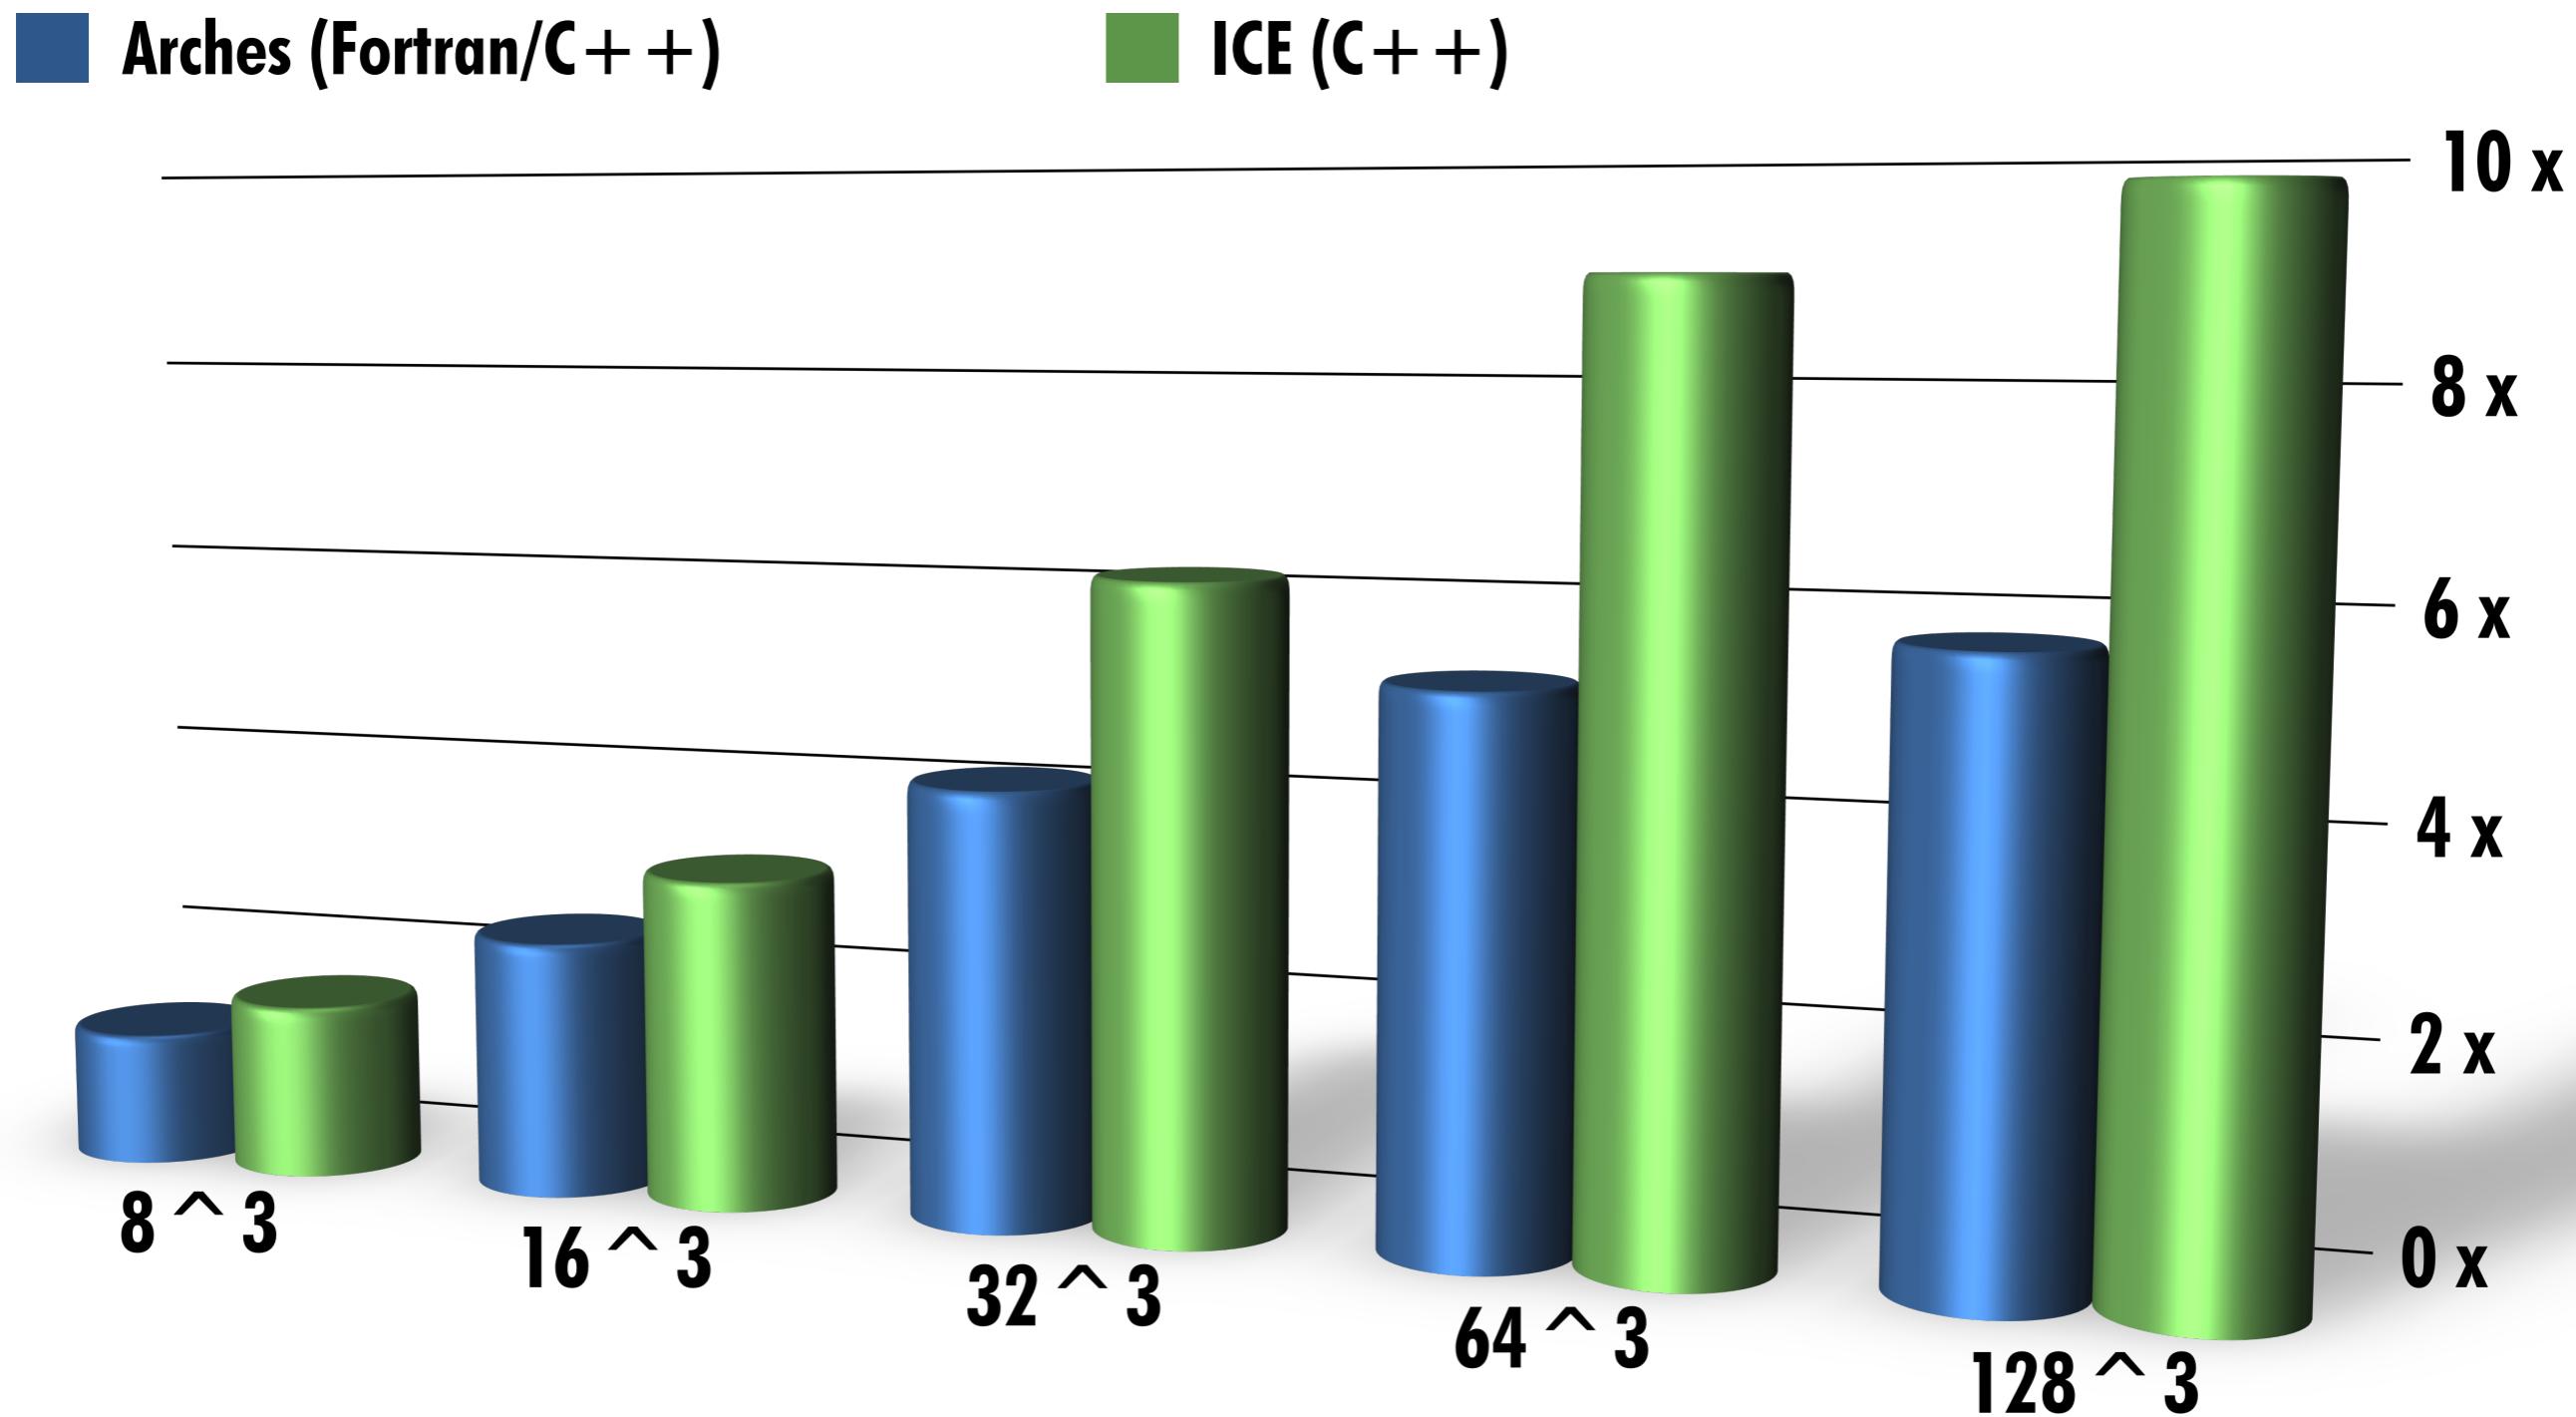
Wasatch

Task Interface

schedule ExprLib
graph as a Uintah
task

Uintah

Speedup using Nebo DSL Syntax compared to traditional iterative/ijk loops in Uintah



In Practice

Solve

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla \phi)$$

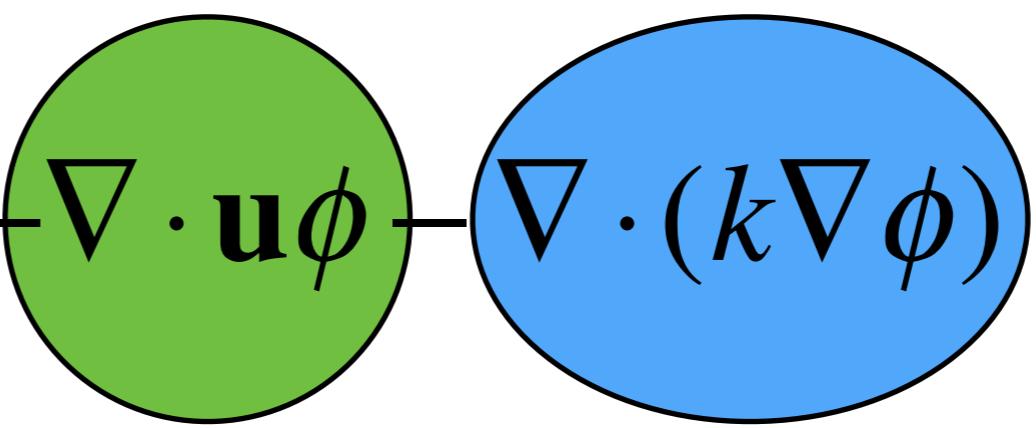
$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla \phi)$$

1. Define expressions (easy)

1. Define expressions (easy)

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u} \phi - \nabla \cdot (k \nabla \phi)$$

Convective Flux Diffusive Flux



$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla \phi)$$

2. Define input spec (easy)

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla\phi)$$

2. Define input spec (easy)

```
<TransportEquation equation="generic">
    <SolutionVariable>phi</SolutionVariable>

    <DiffusiveFlux direction="X" coefficient="0.4"/>
    <DiffusiveFlux direction="Y" coefficient="0.2"/>
    <DiffusiveFlux direction="Z" coefficient="0.1"/>

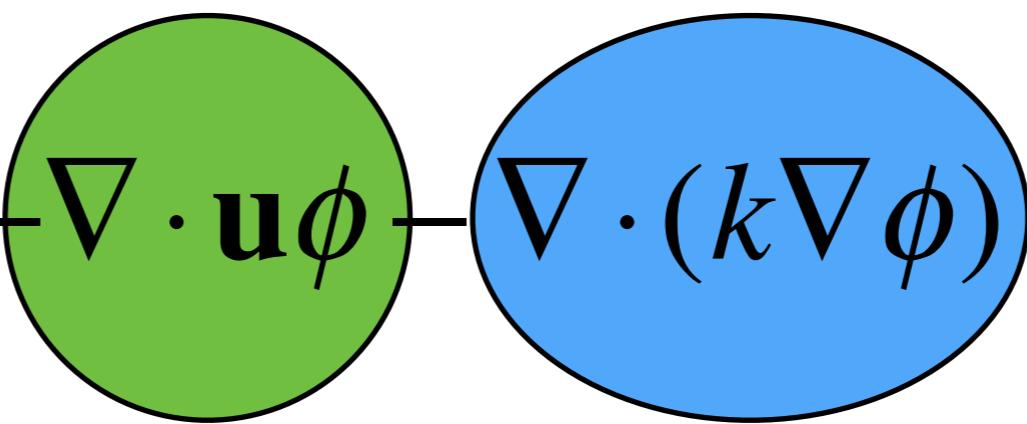
    <ConvectiveFlux direction="Y" method="SUPERBEE">
        <AdvectiveVelocity>
            <NameTag name="v"/>
        </AdvectiveVelocity>
    </ConvectiveFlux>

</TransportEquation>
```

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla \phi)$$

3. Implement expressions (easy)

3. Implement expressions (easy)

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla\phi)$$


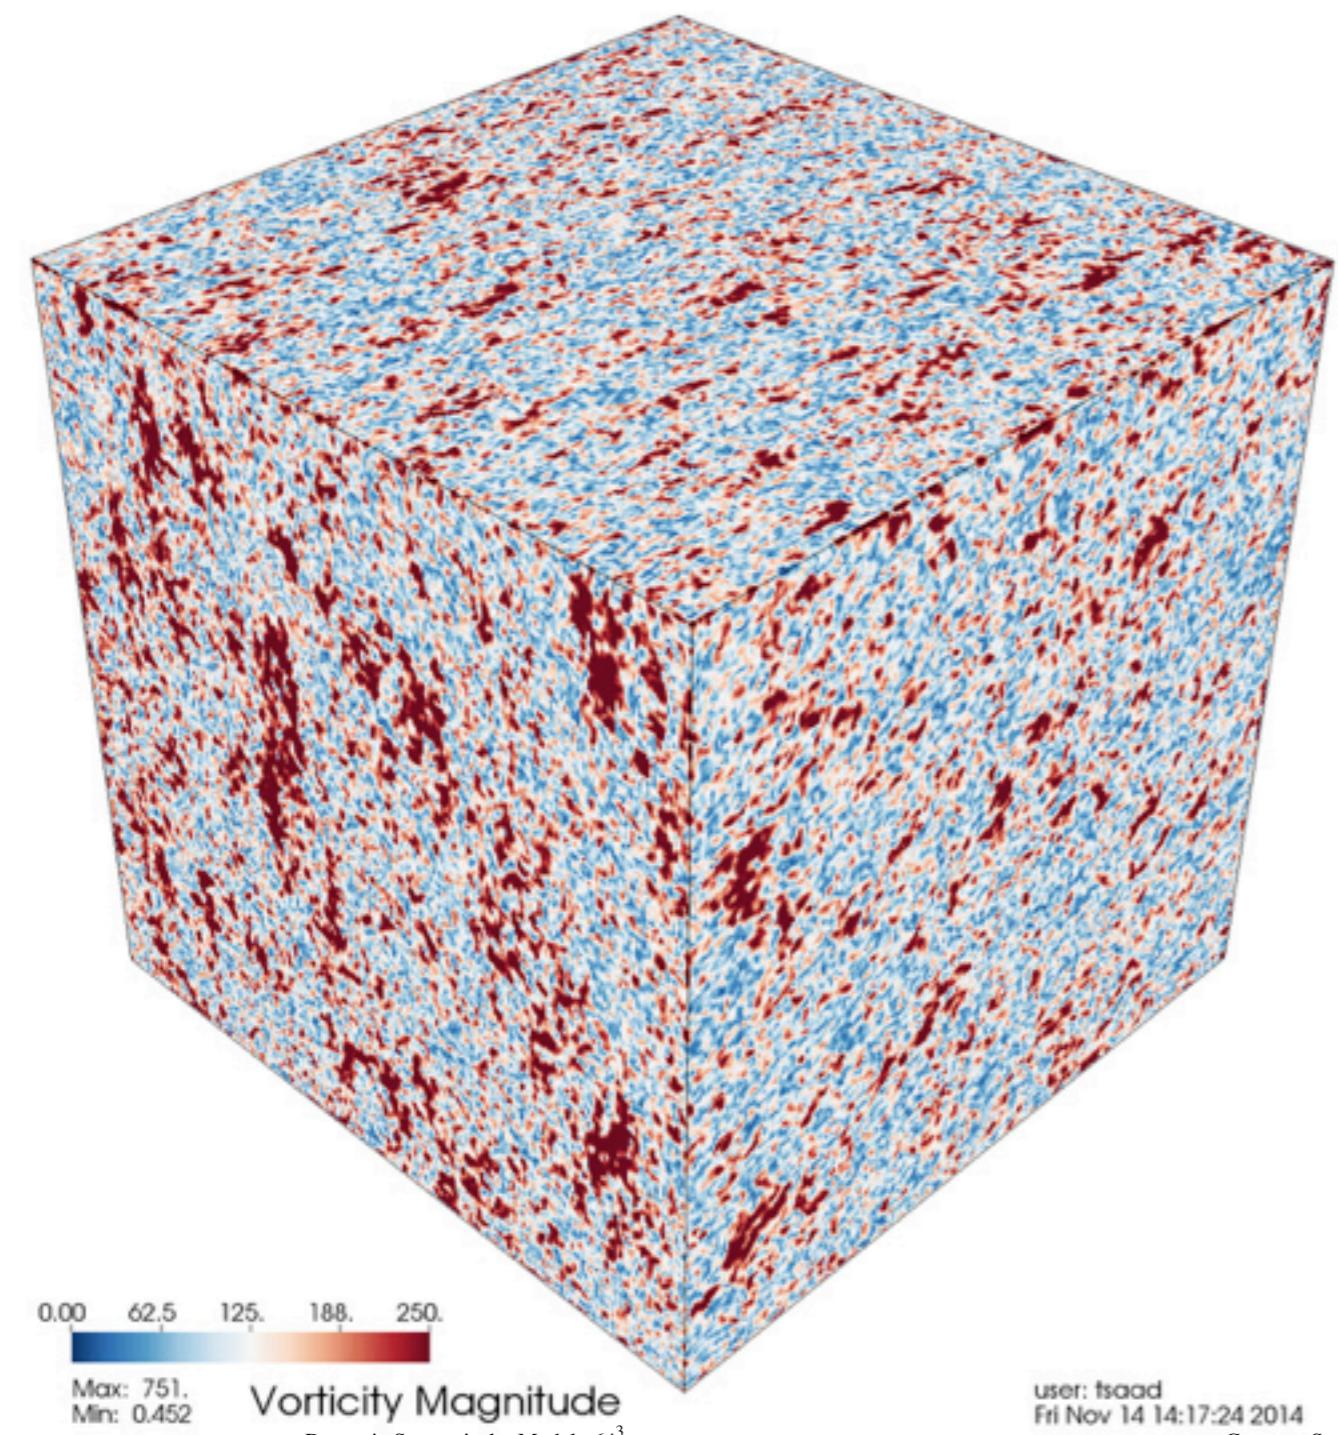
`cflux <= div(u*interp(phi))`

`dflux <= div(interp(k)*grad(phi))`

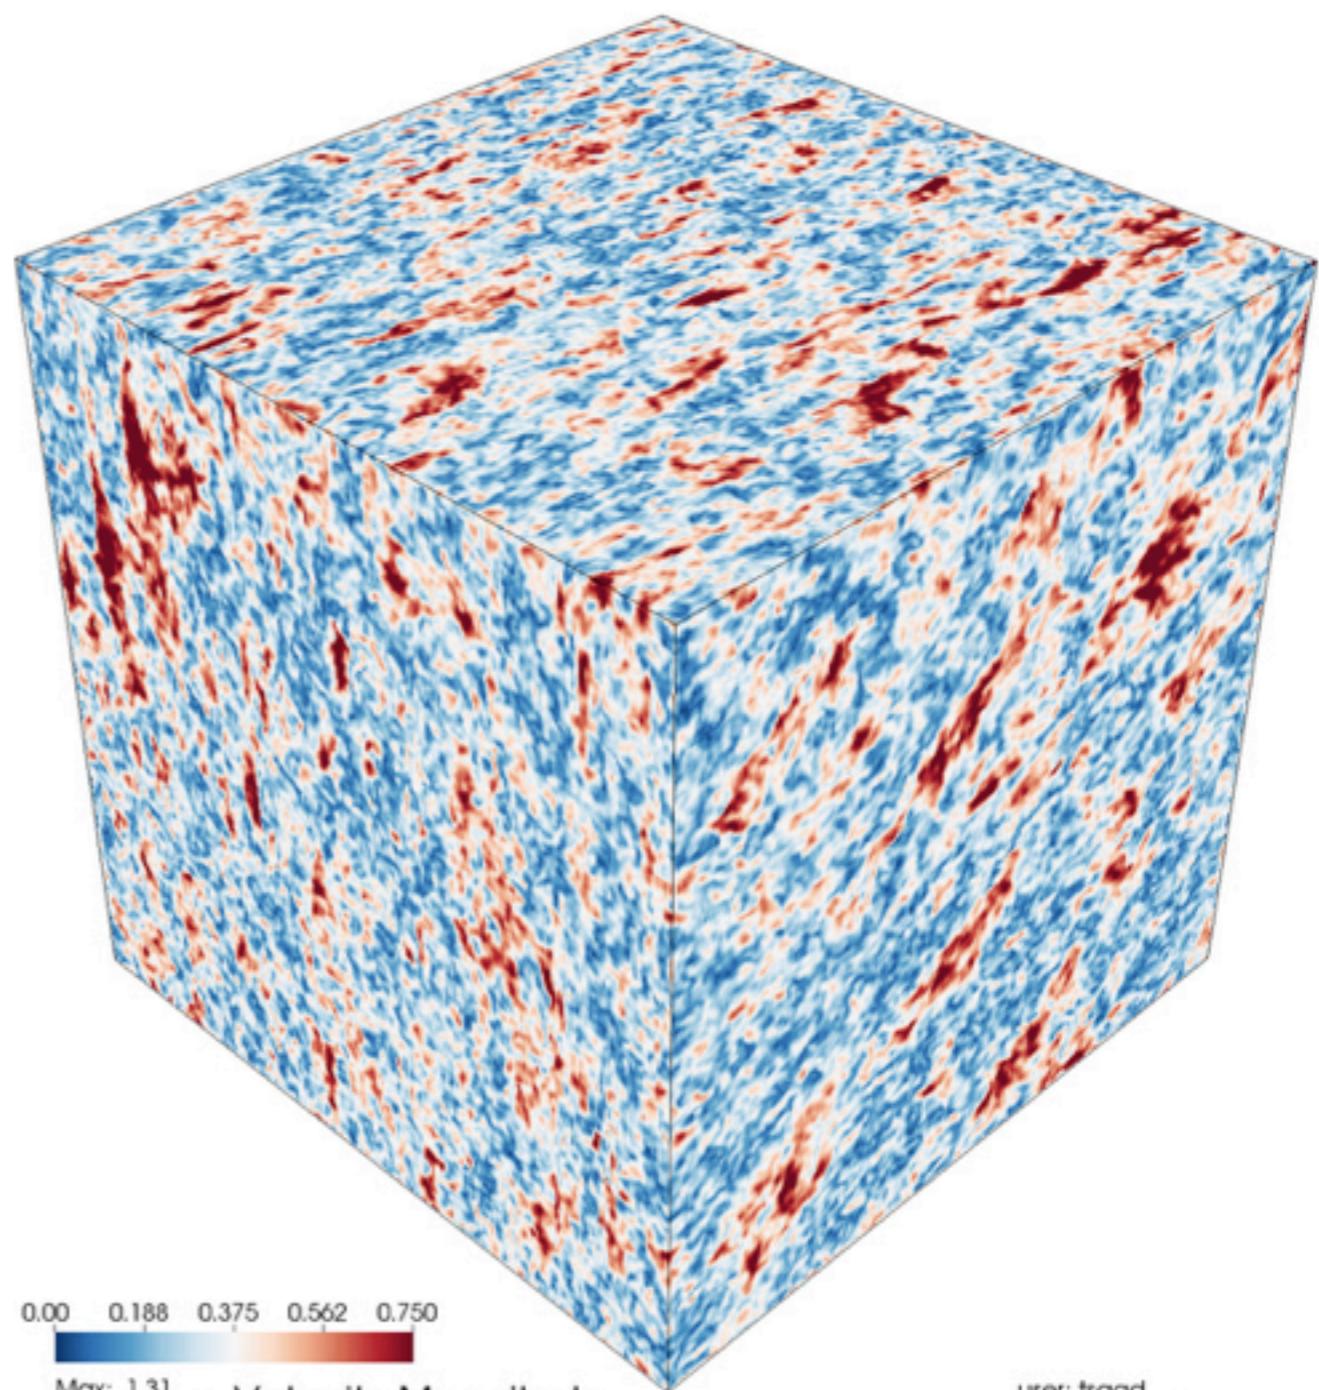
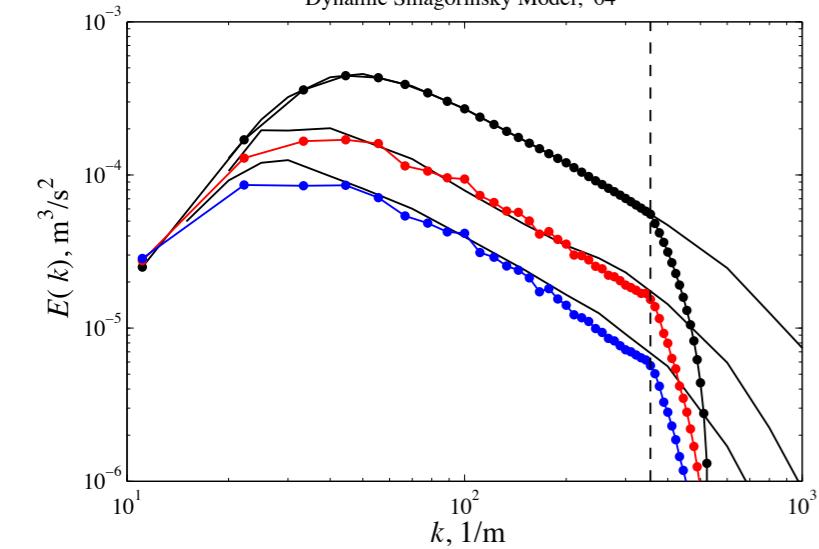
Interaction with computer science

$$\frac{\partial \phi}{\partial t} = -\nabla \cdot \mathbf{u}\phi - \nabla \cdot (k\nabla \phi)$$

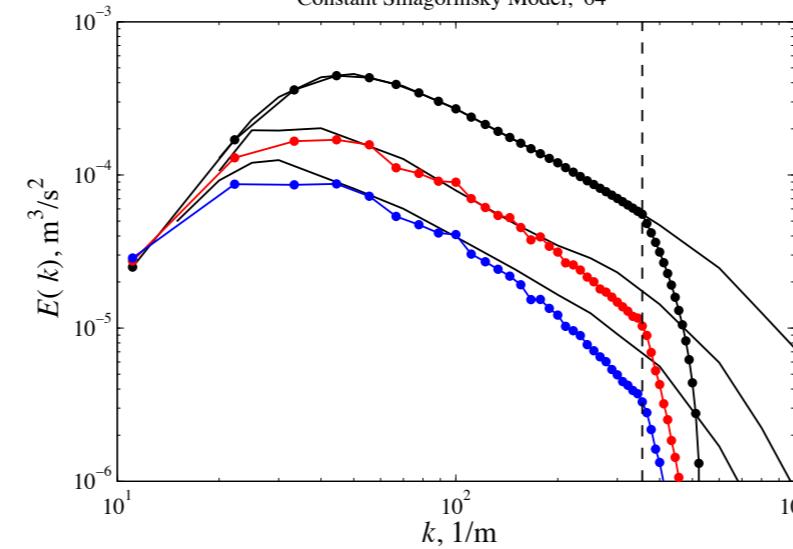
4. Test, verify, validate



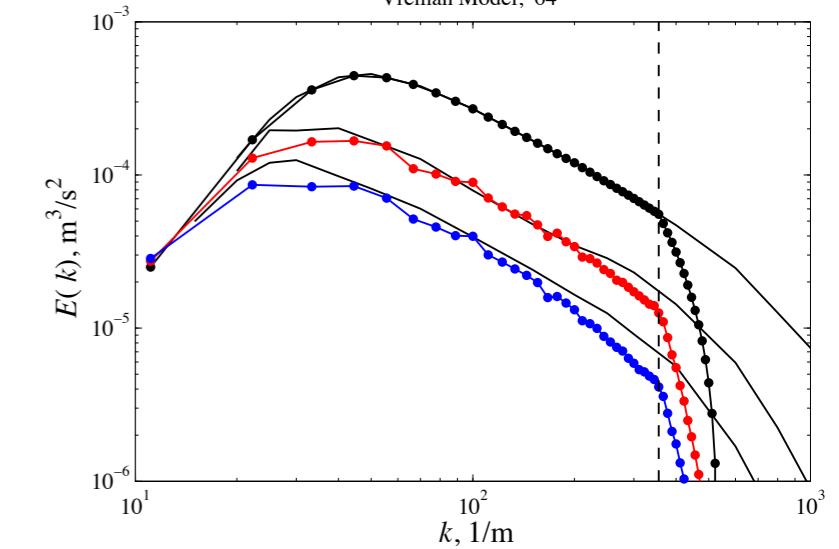
Max: 751.
Min: 0.452
Vorticity Magnitude
Dynamic Smagorinsky Model, 64^3

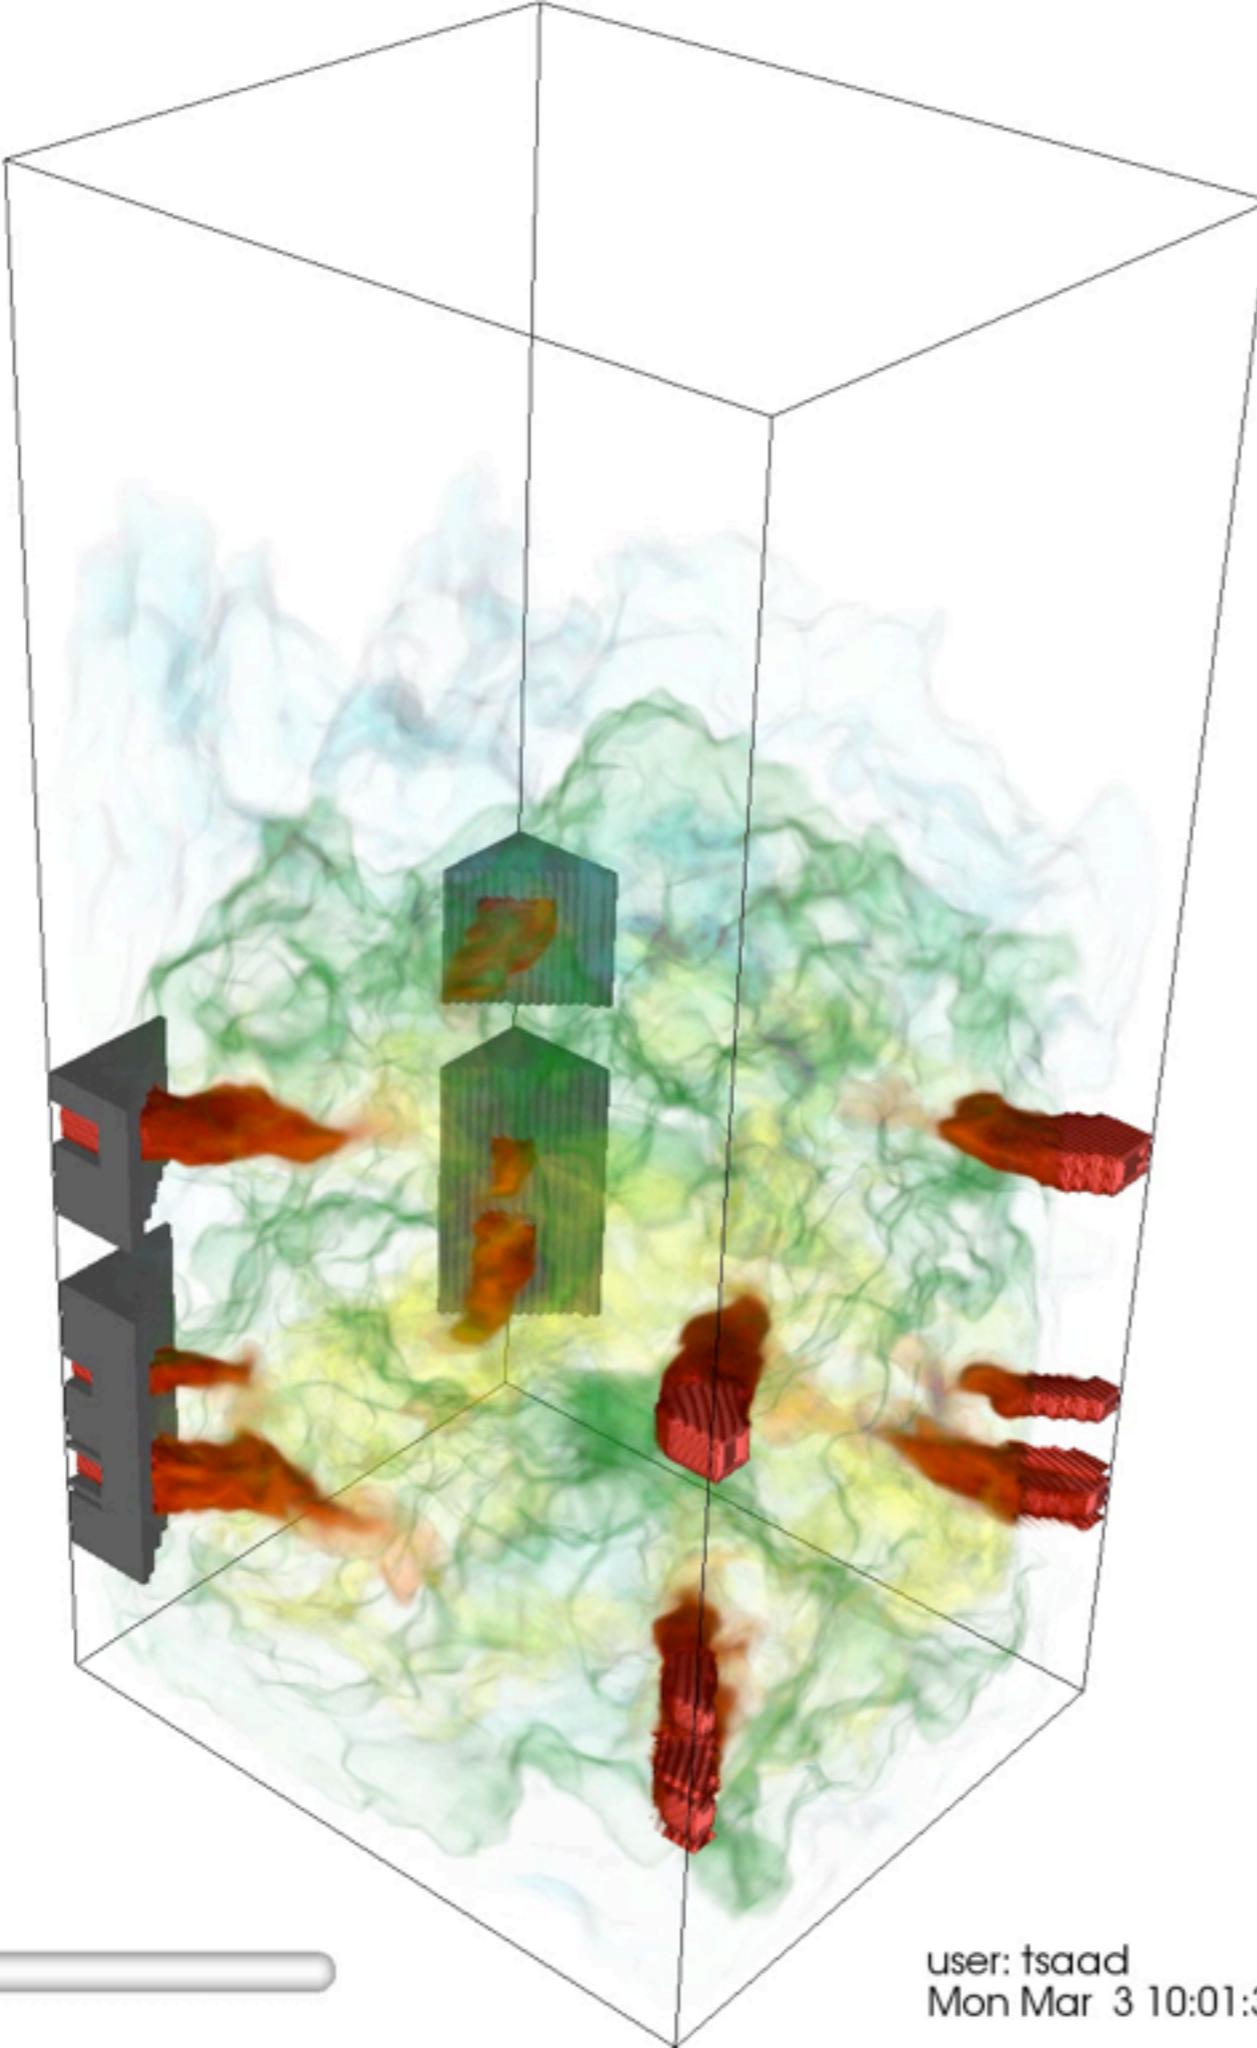
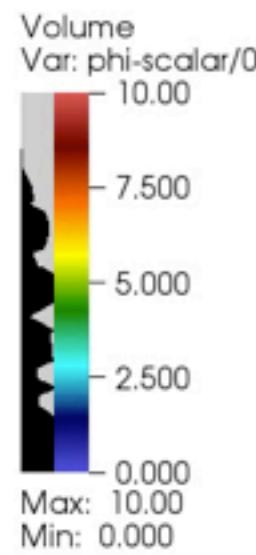


user: tsaad
Fri Nov 14 14:17:24 2014
Max: 1.31
Min: 0.00125
Velocity Magnitude
Constant Smagorinsky Model, 64^3



user: tsaad
Fri Nov 14 15:00:24 2014
Vreman Model, 64^3

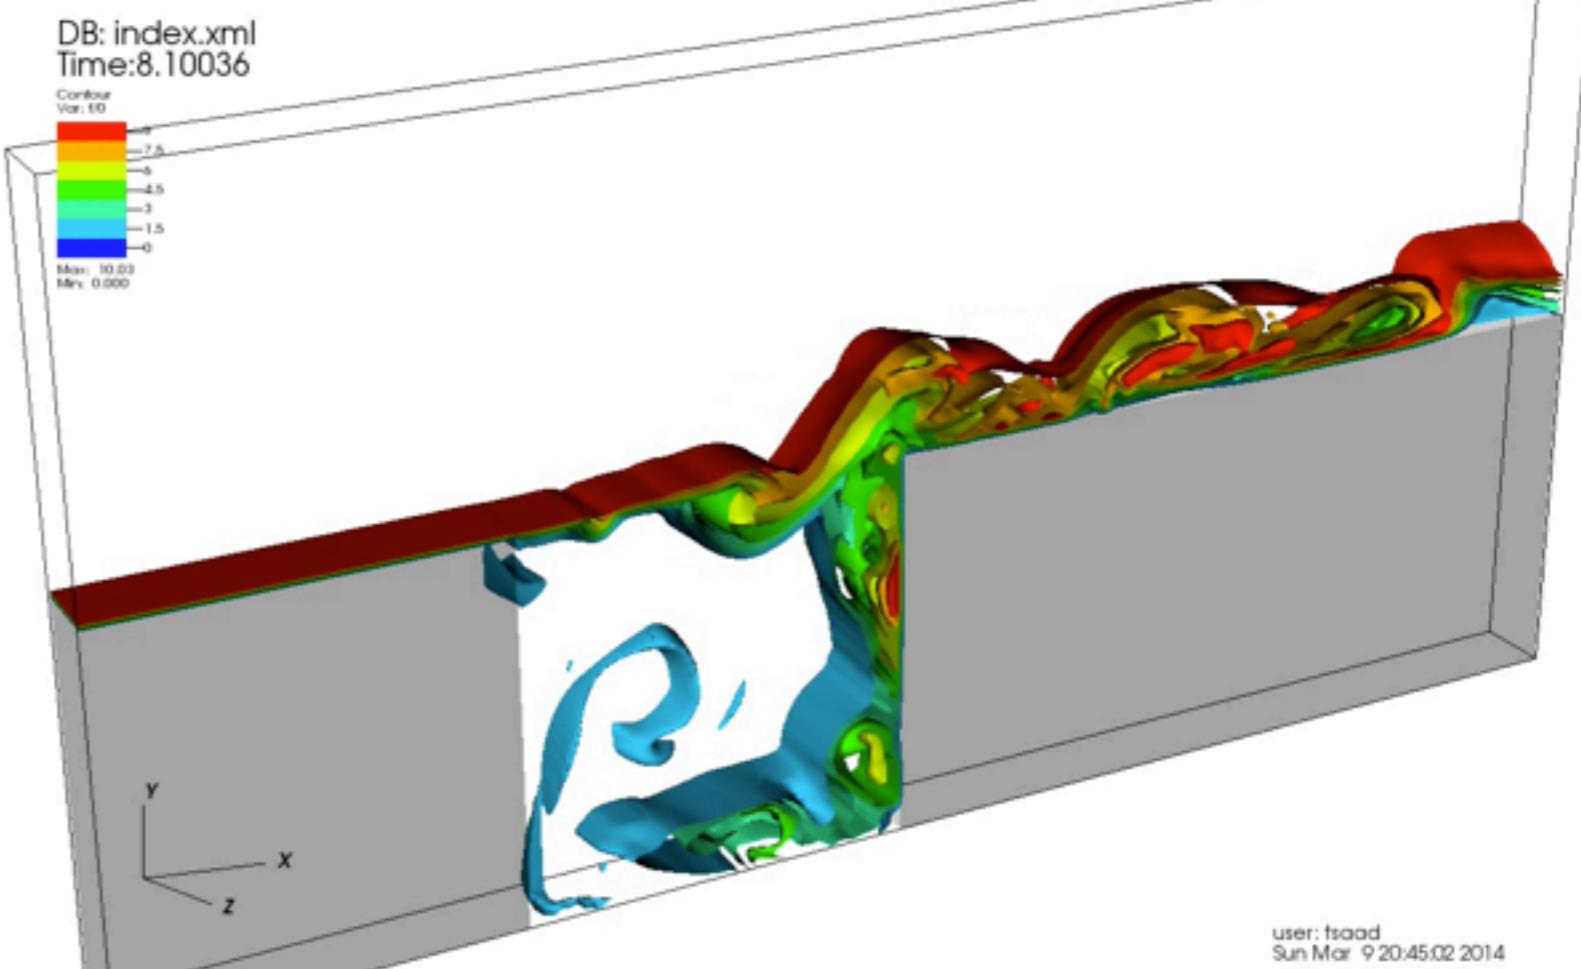
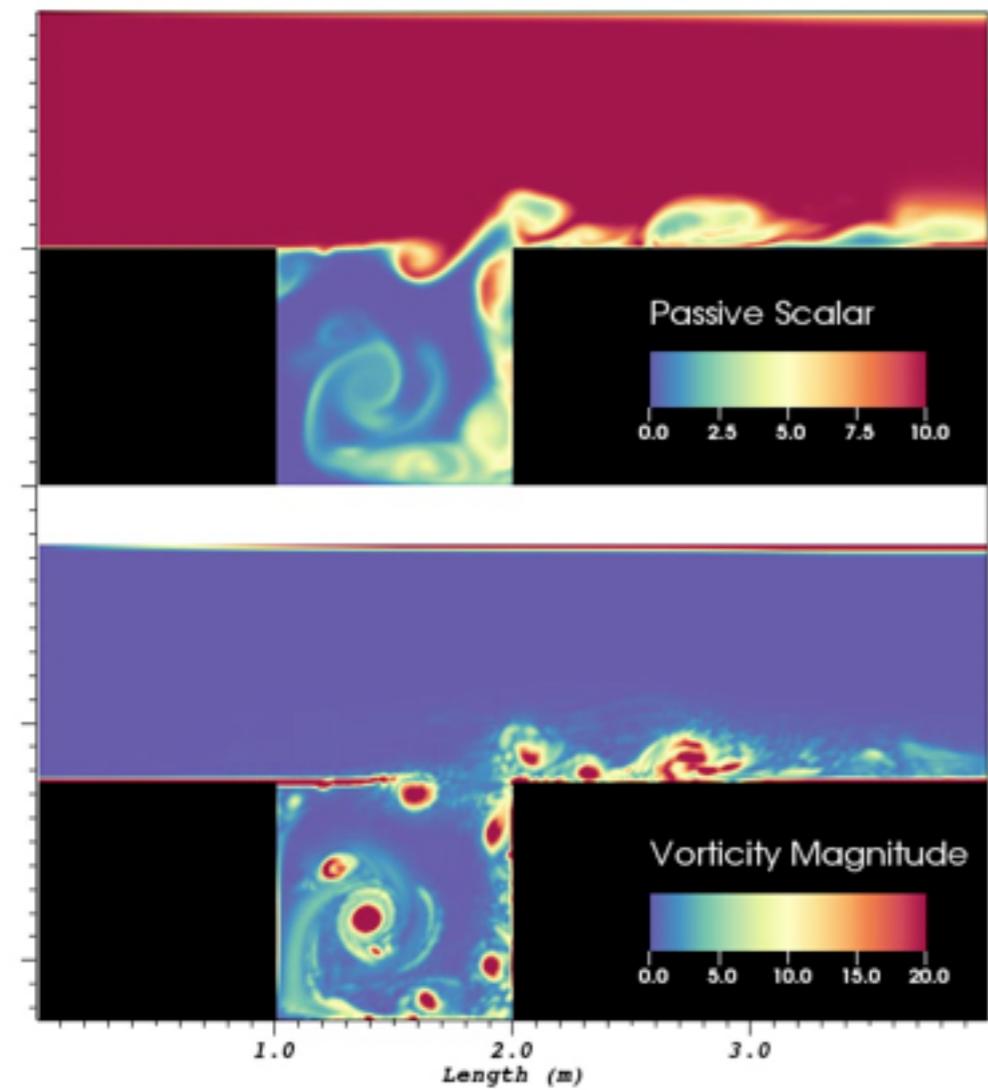




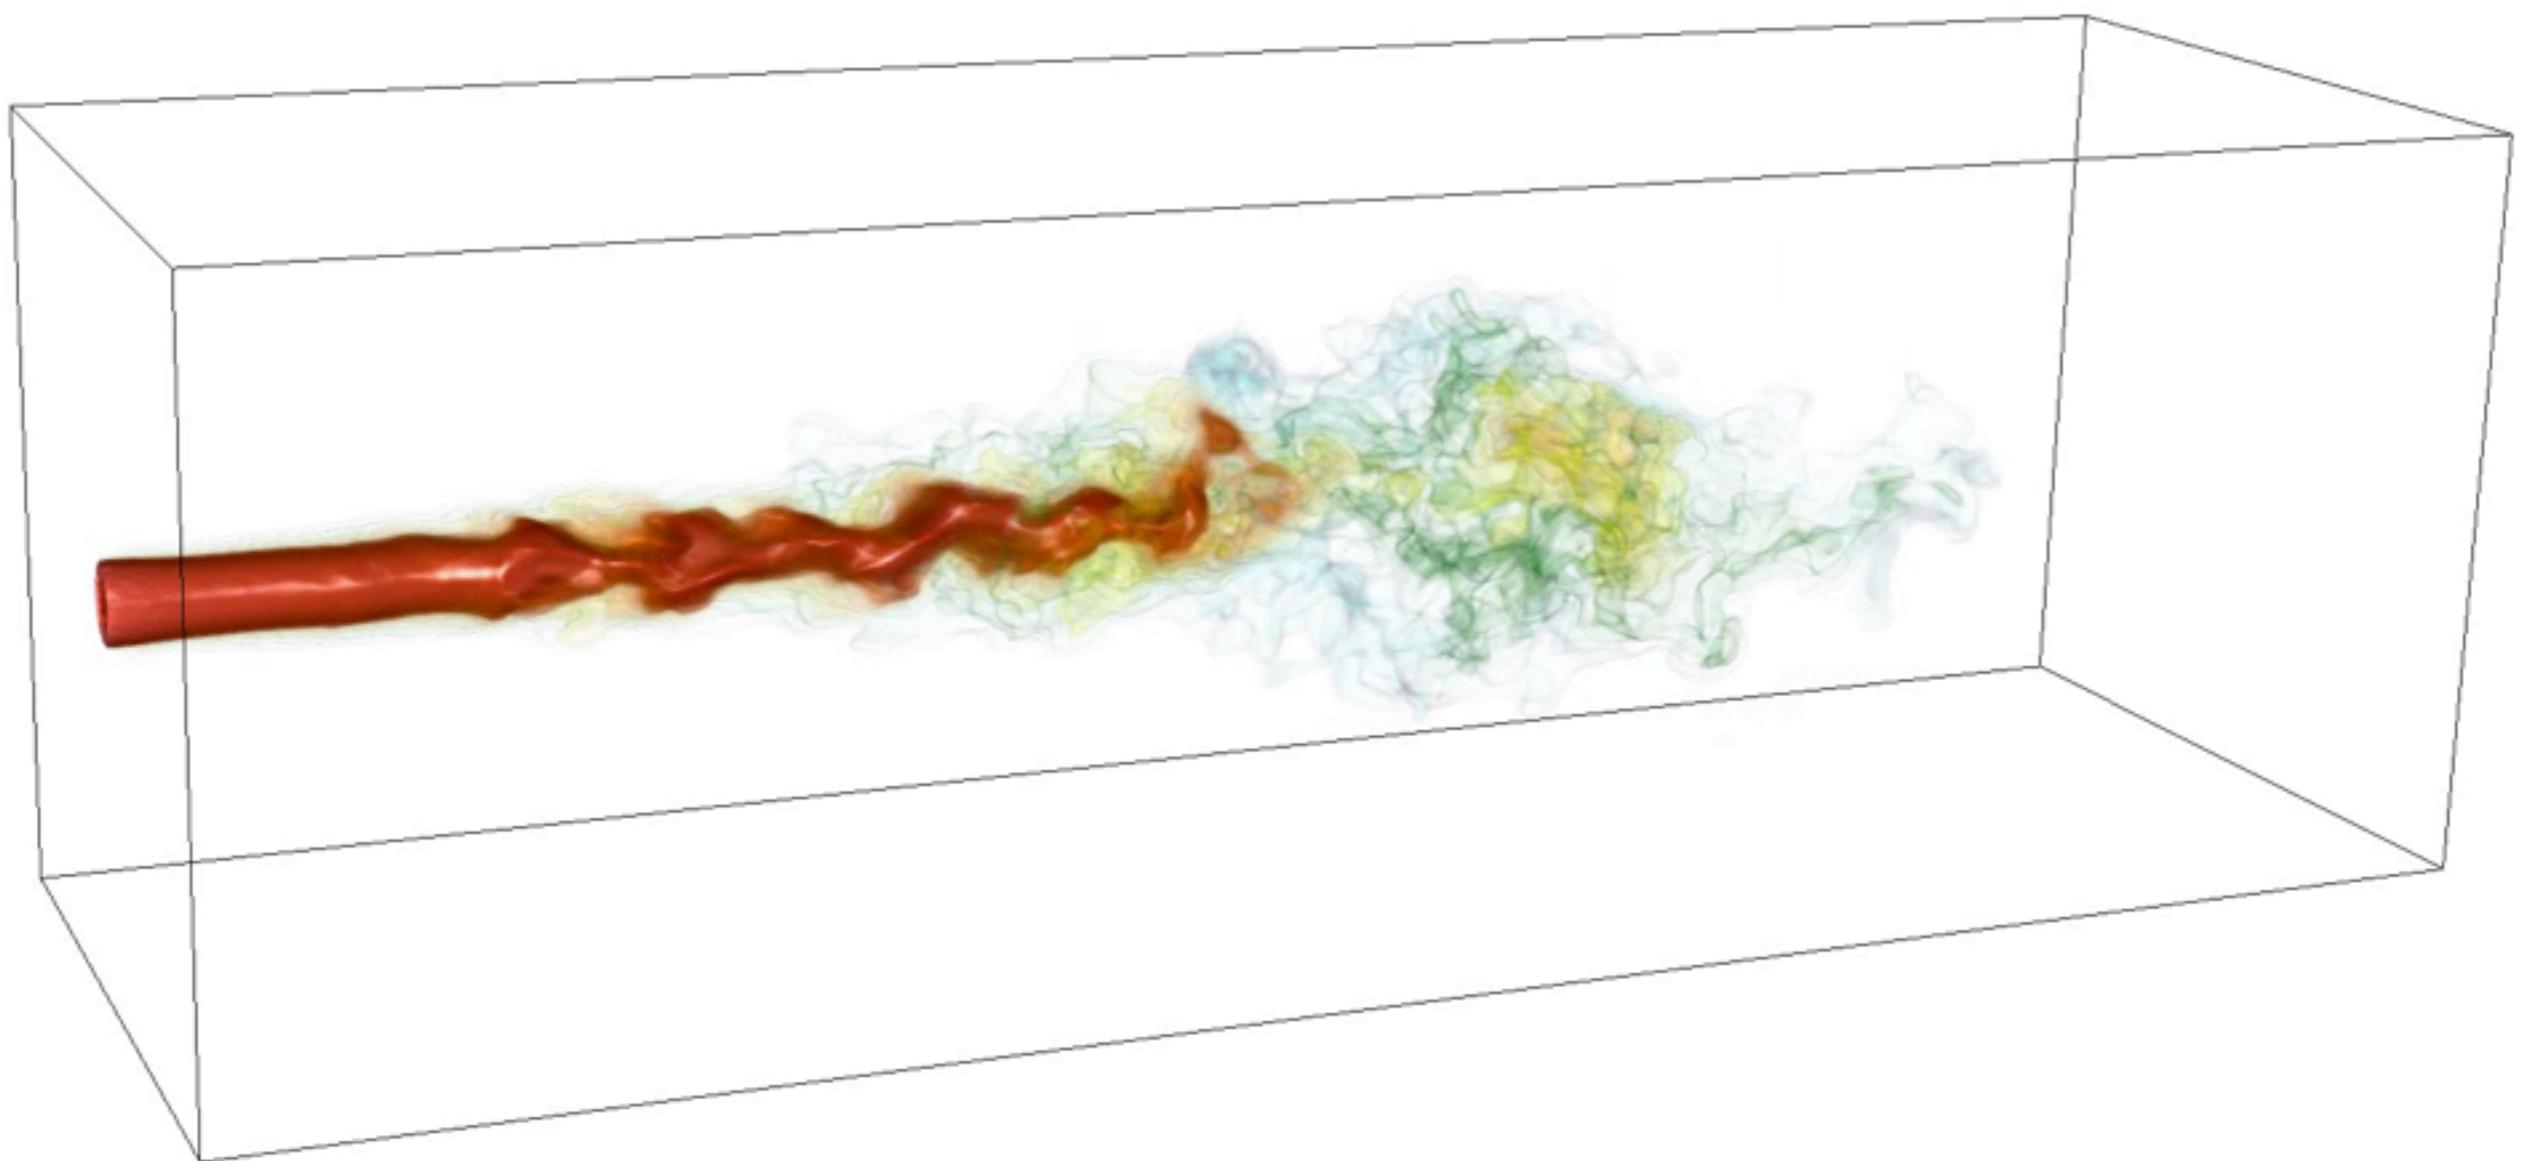
Time=0.000 s

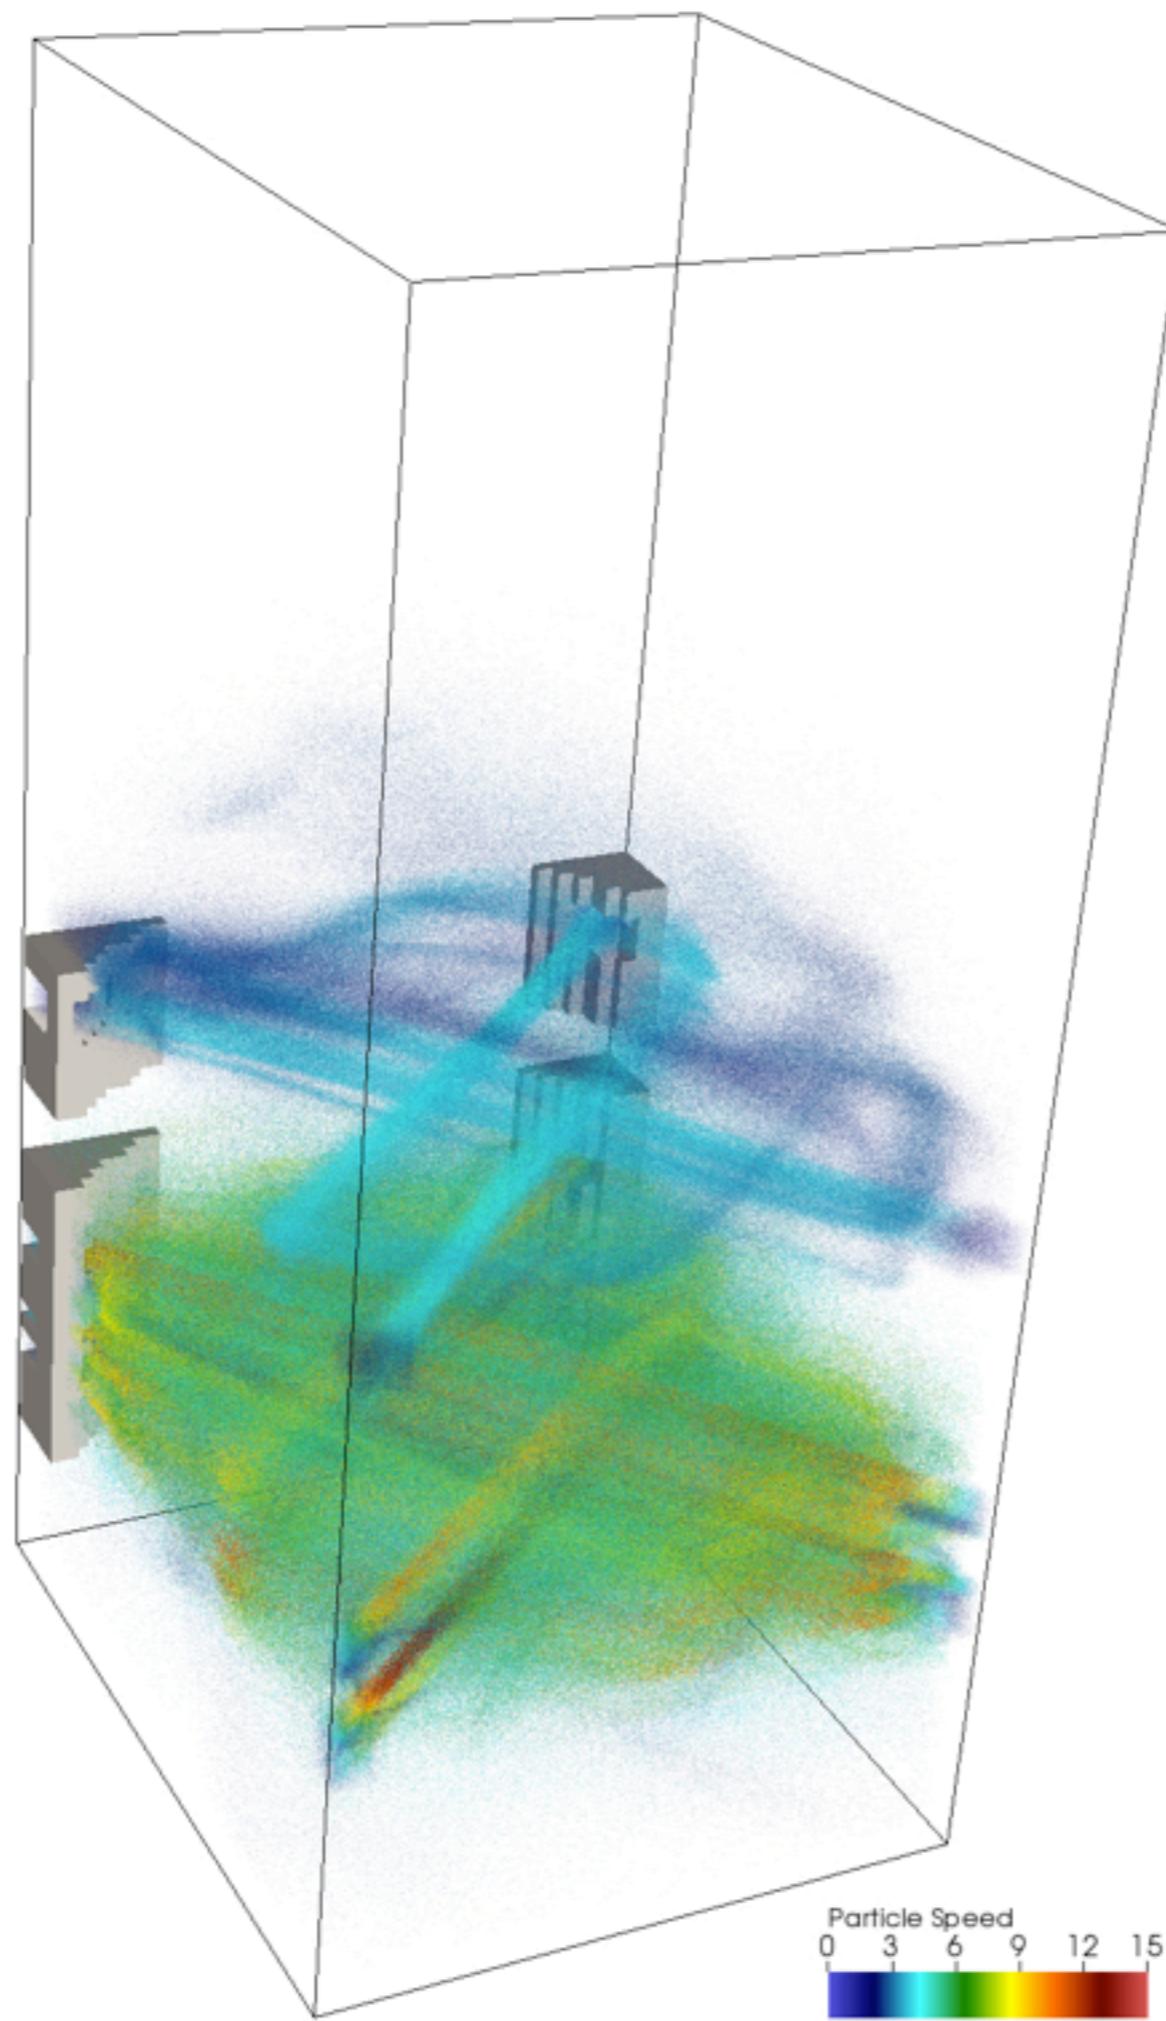
user: tsaad
Mon Mar 3 10:14:12 2014

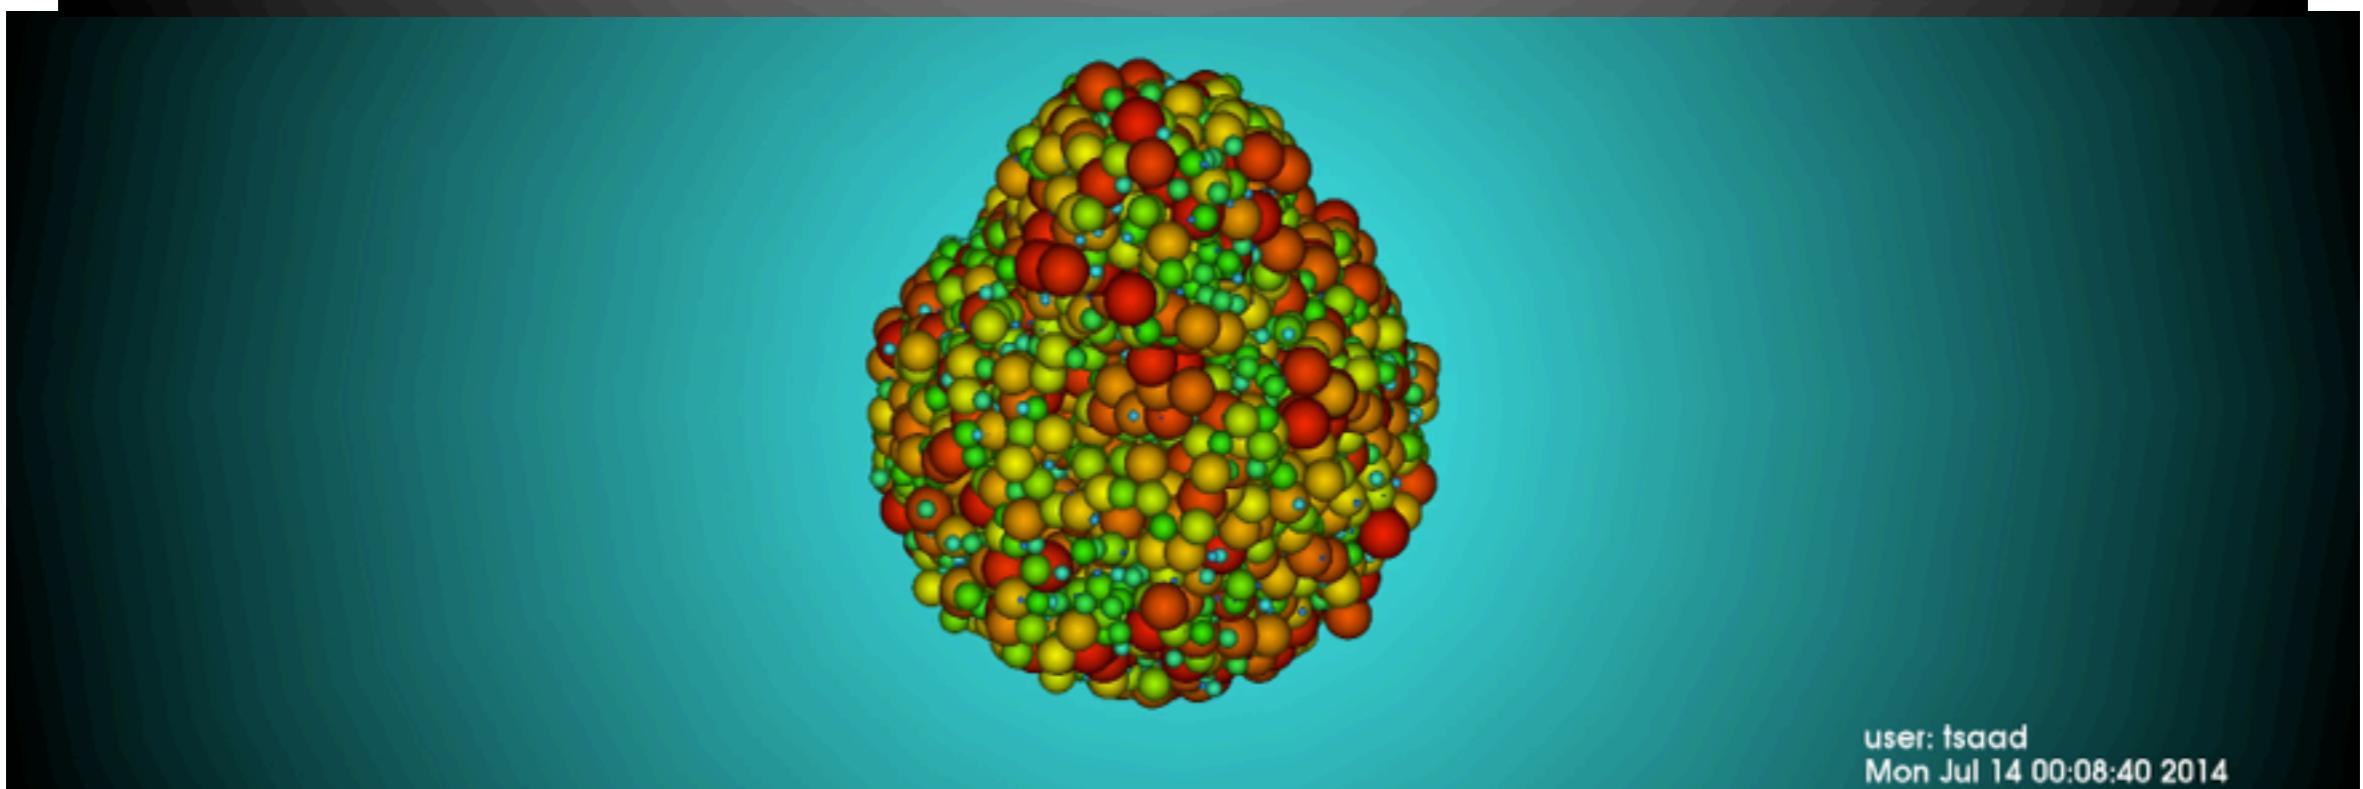
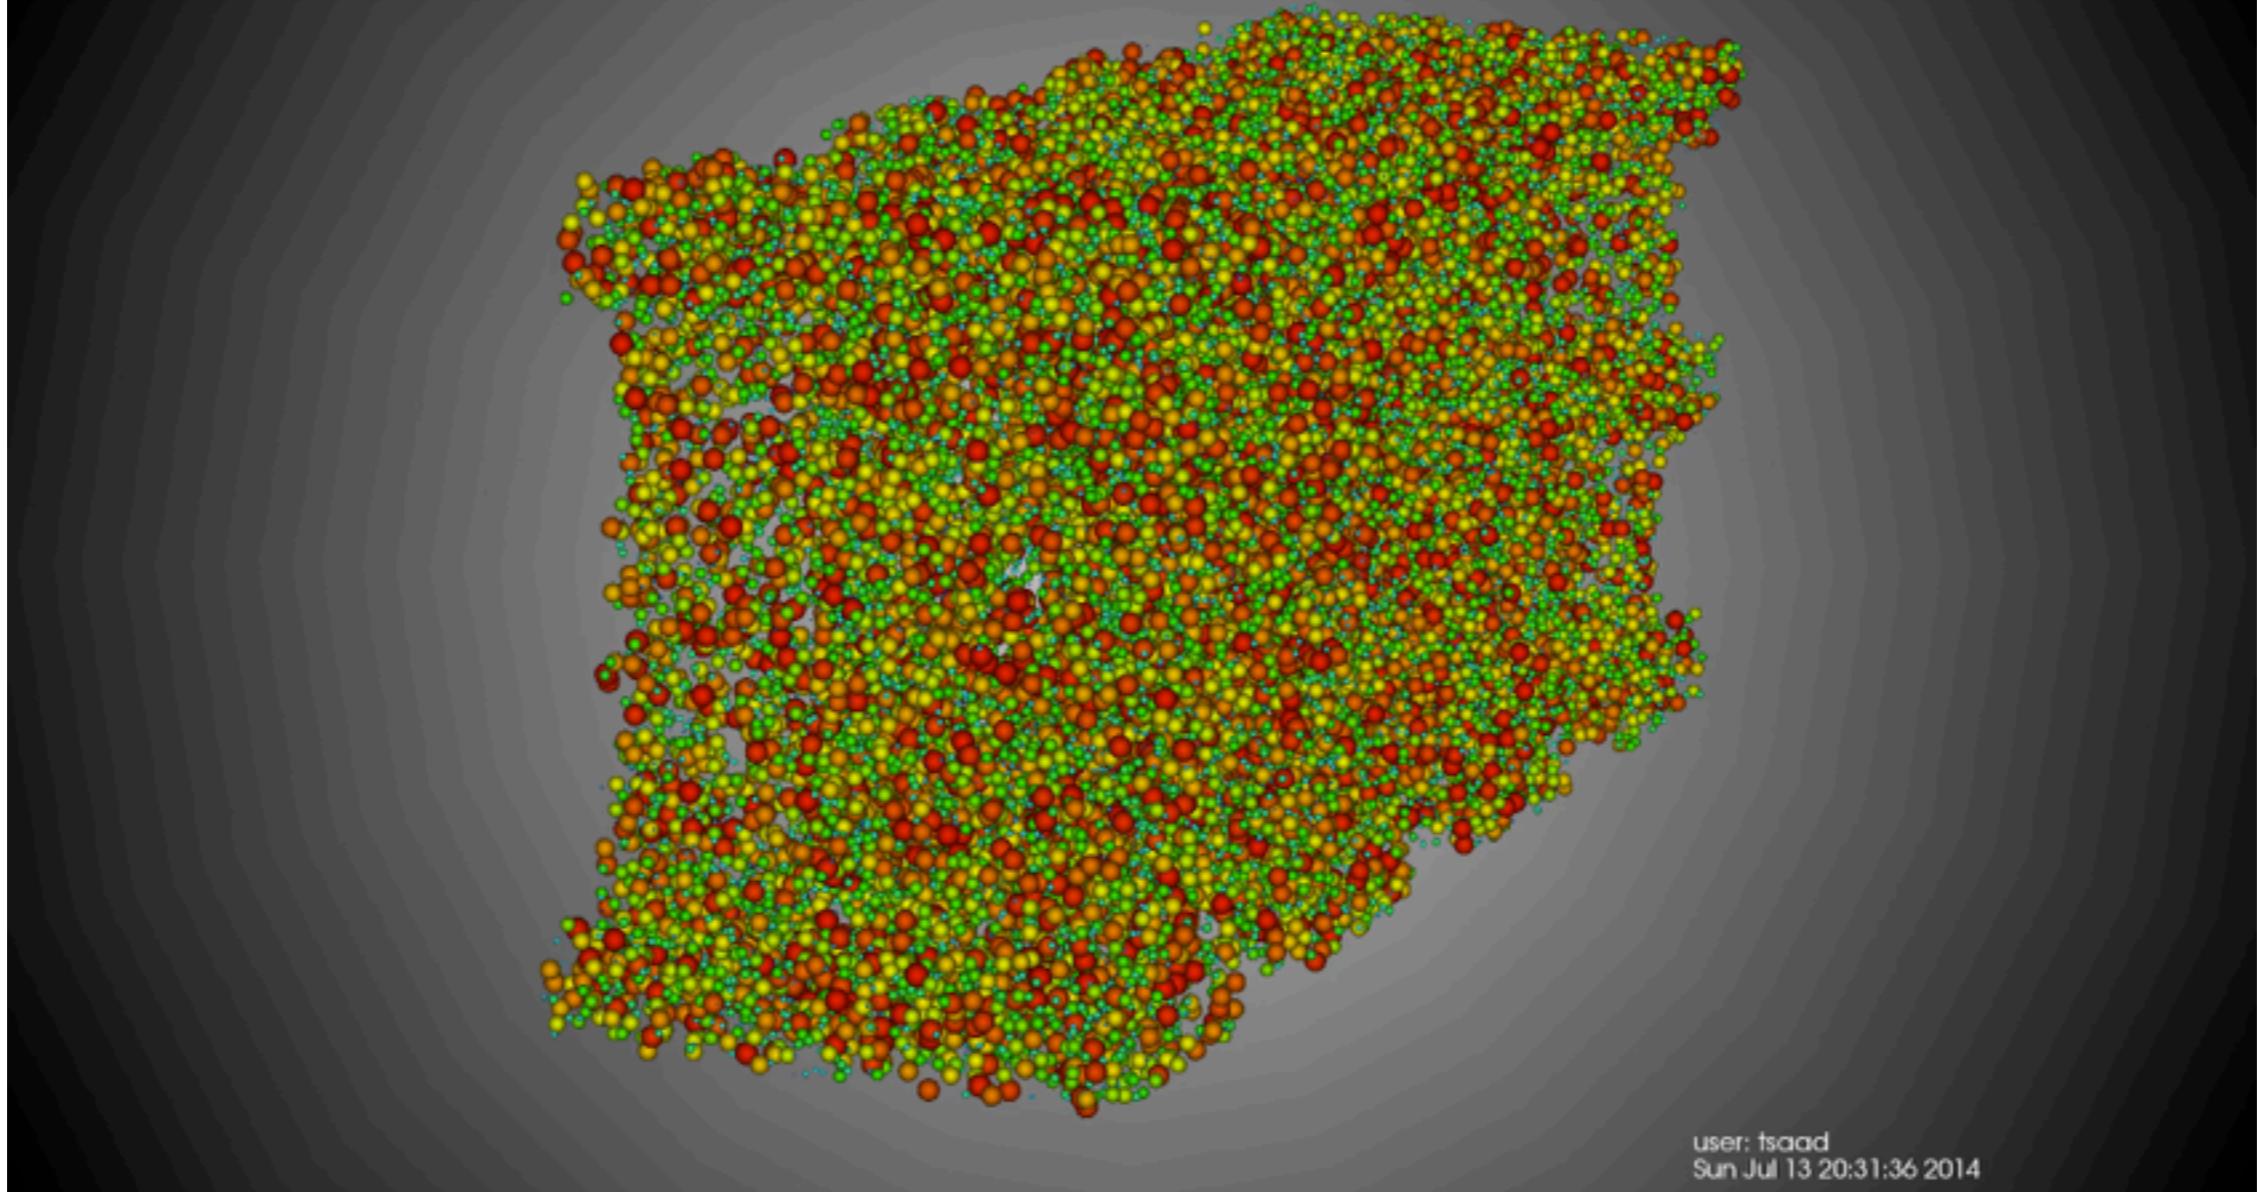
Time=0.781 s
user: tsaad
Mon Mar 3 10:01:39 201



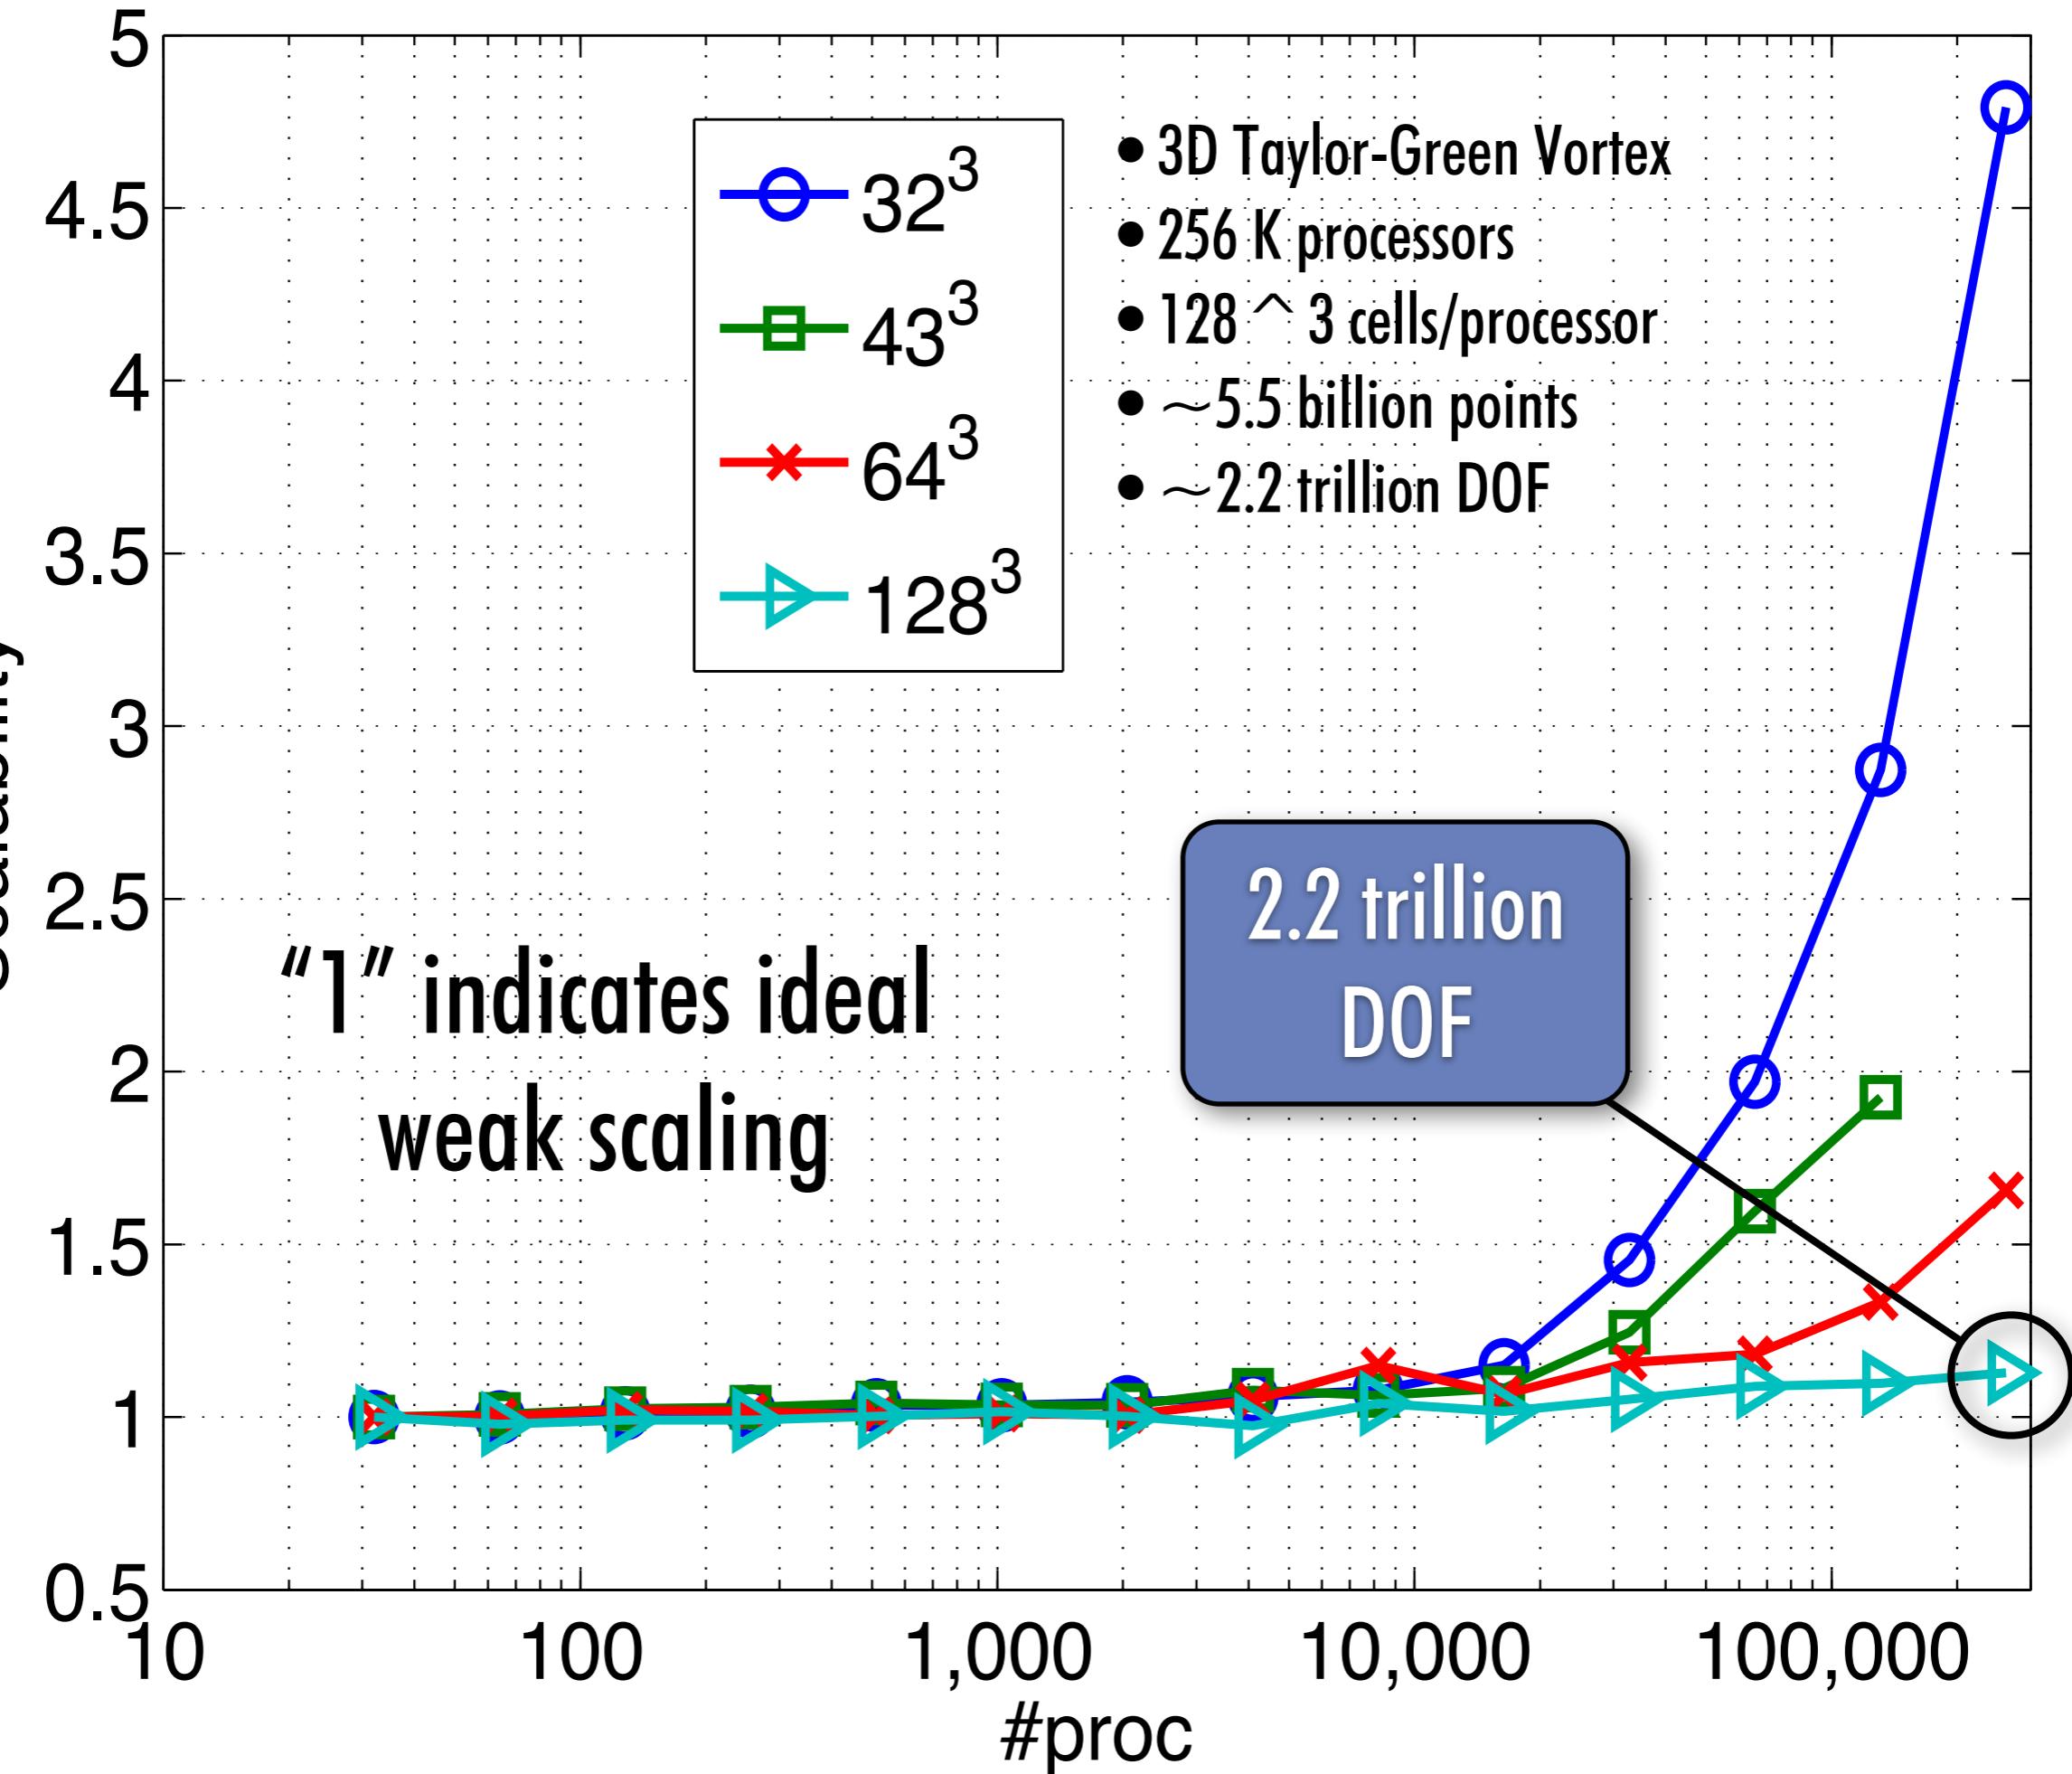
user: tsaad
Sun Mar 9 20:39:13 20







Weak scaling on Titan

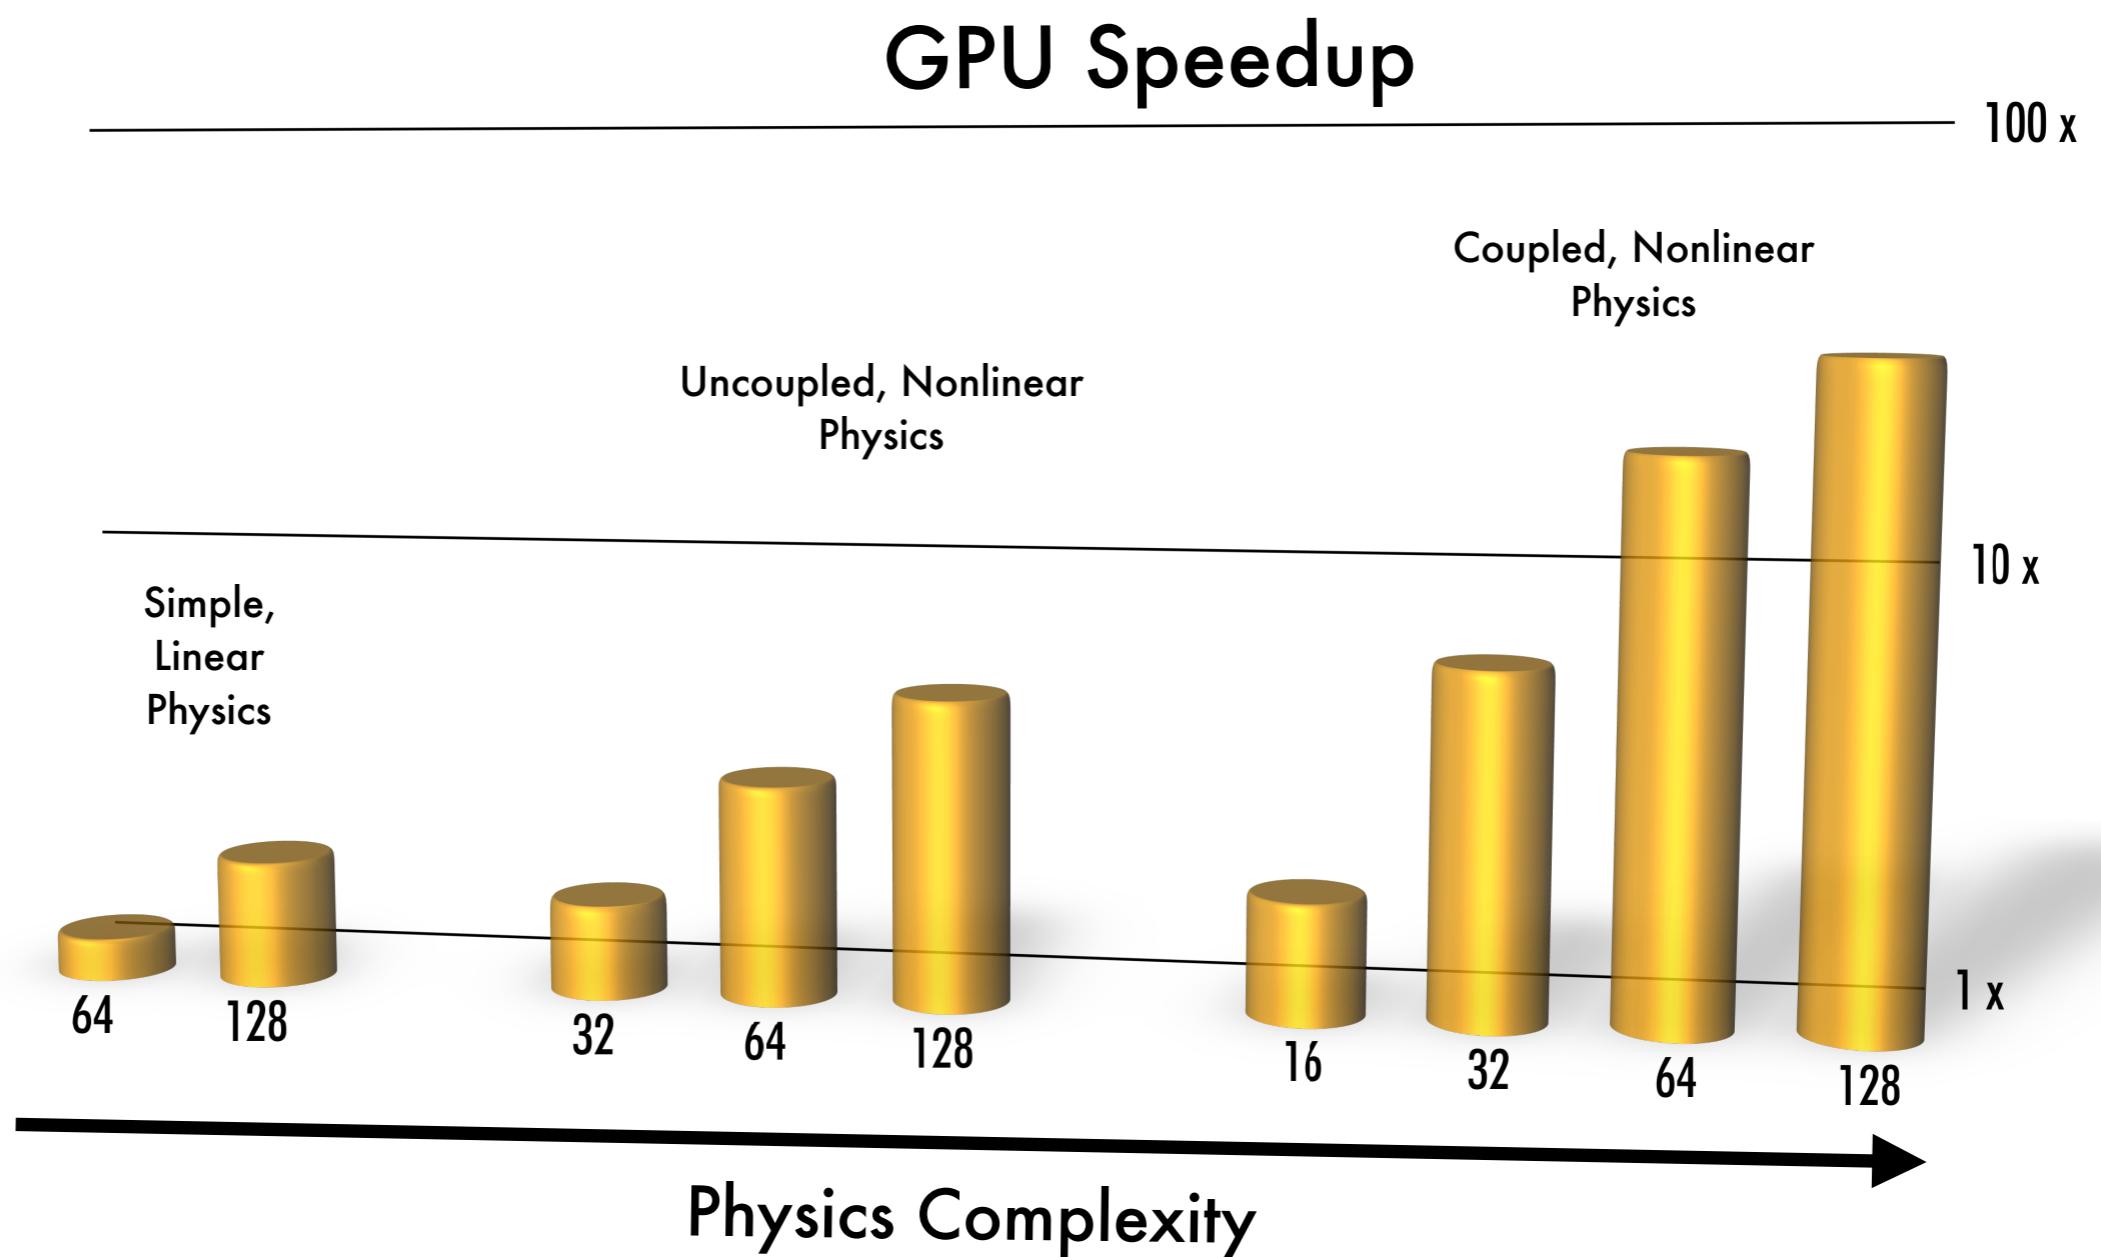


GPU Scaling

$$\frac{\partial \phi_i}{\partial t} + \nabla \cdot \mathbf{u} \phi_i = -\nabla \cdot \Gamma \nabla \phi_i + S; \quad i = 0, 1, \dots$$

$$S = \begin{cases} 0 & \text{Simple physics} \\ e^{\phi_i} & \text{Moderate} \\ \sum_j e^{\phi_j} & \text{Coupled} \end{cases}$$

Single-Node GPU Speedup



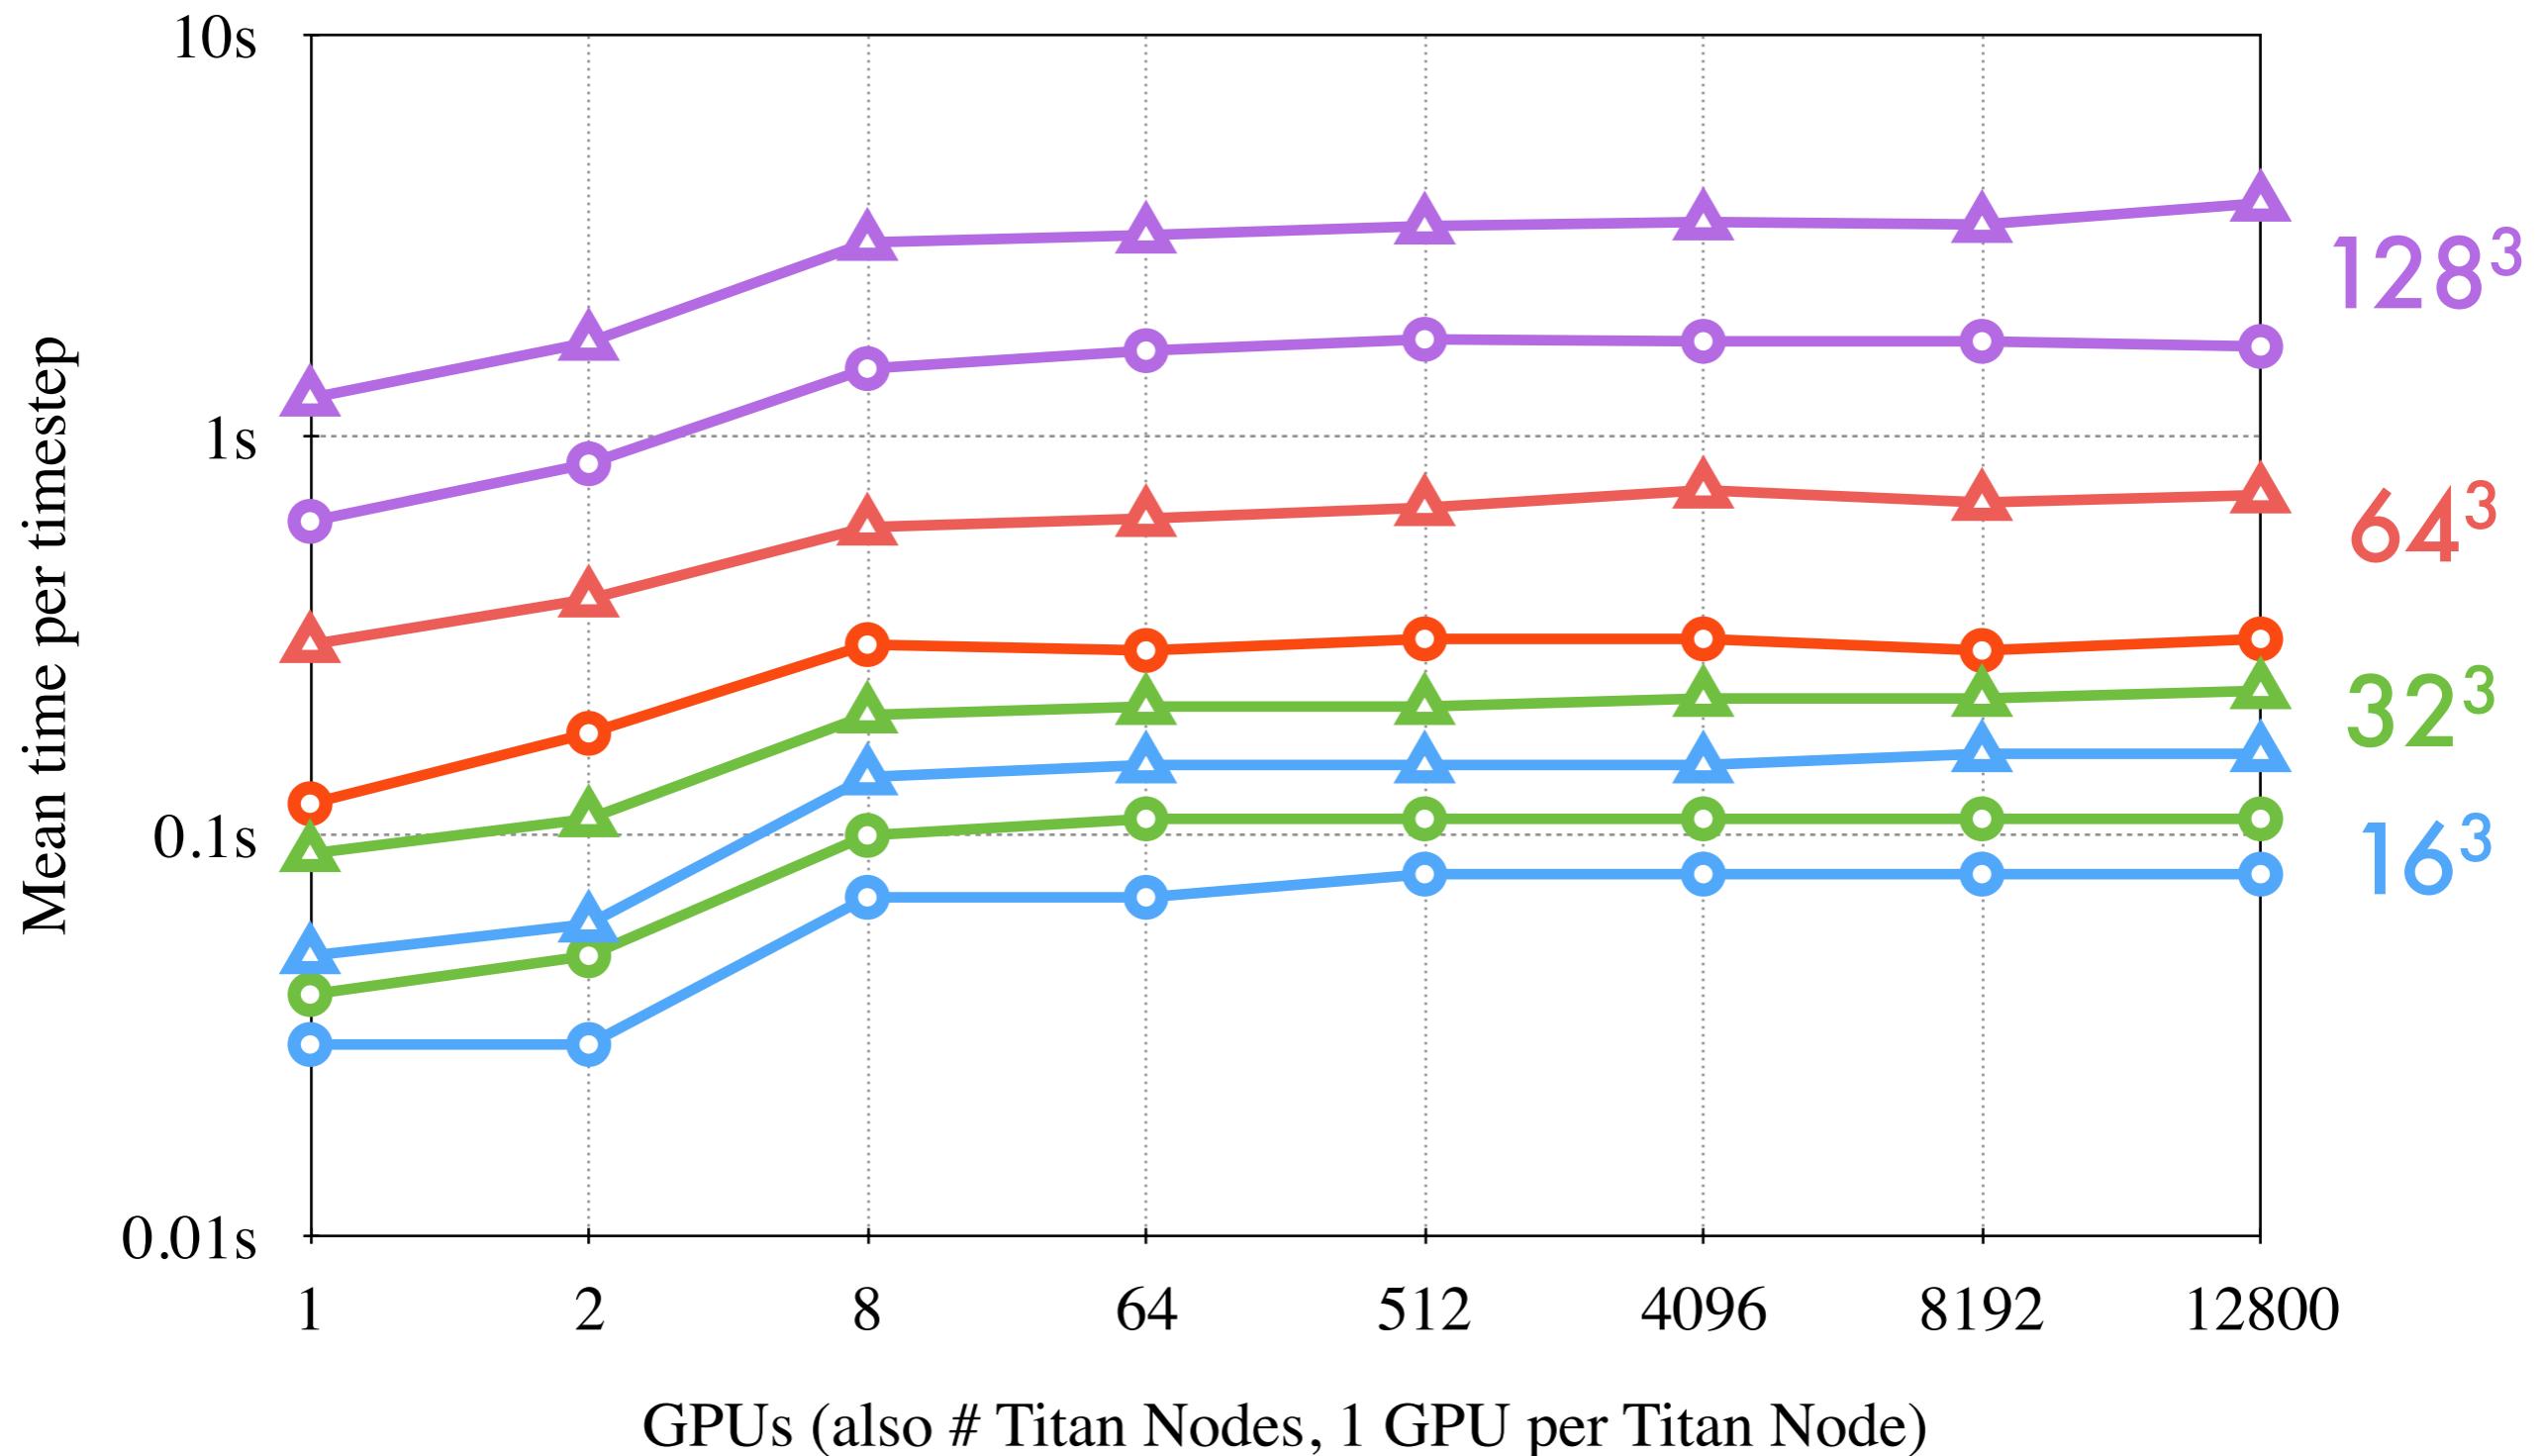
Scaling on Titan

Simple Physics (Scalar Transport without Source Terms)

○ 10 Equations

△ 20 Equations

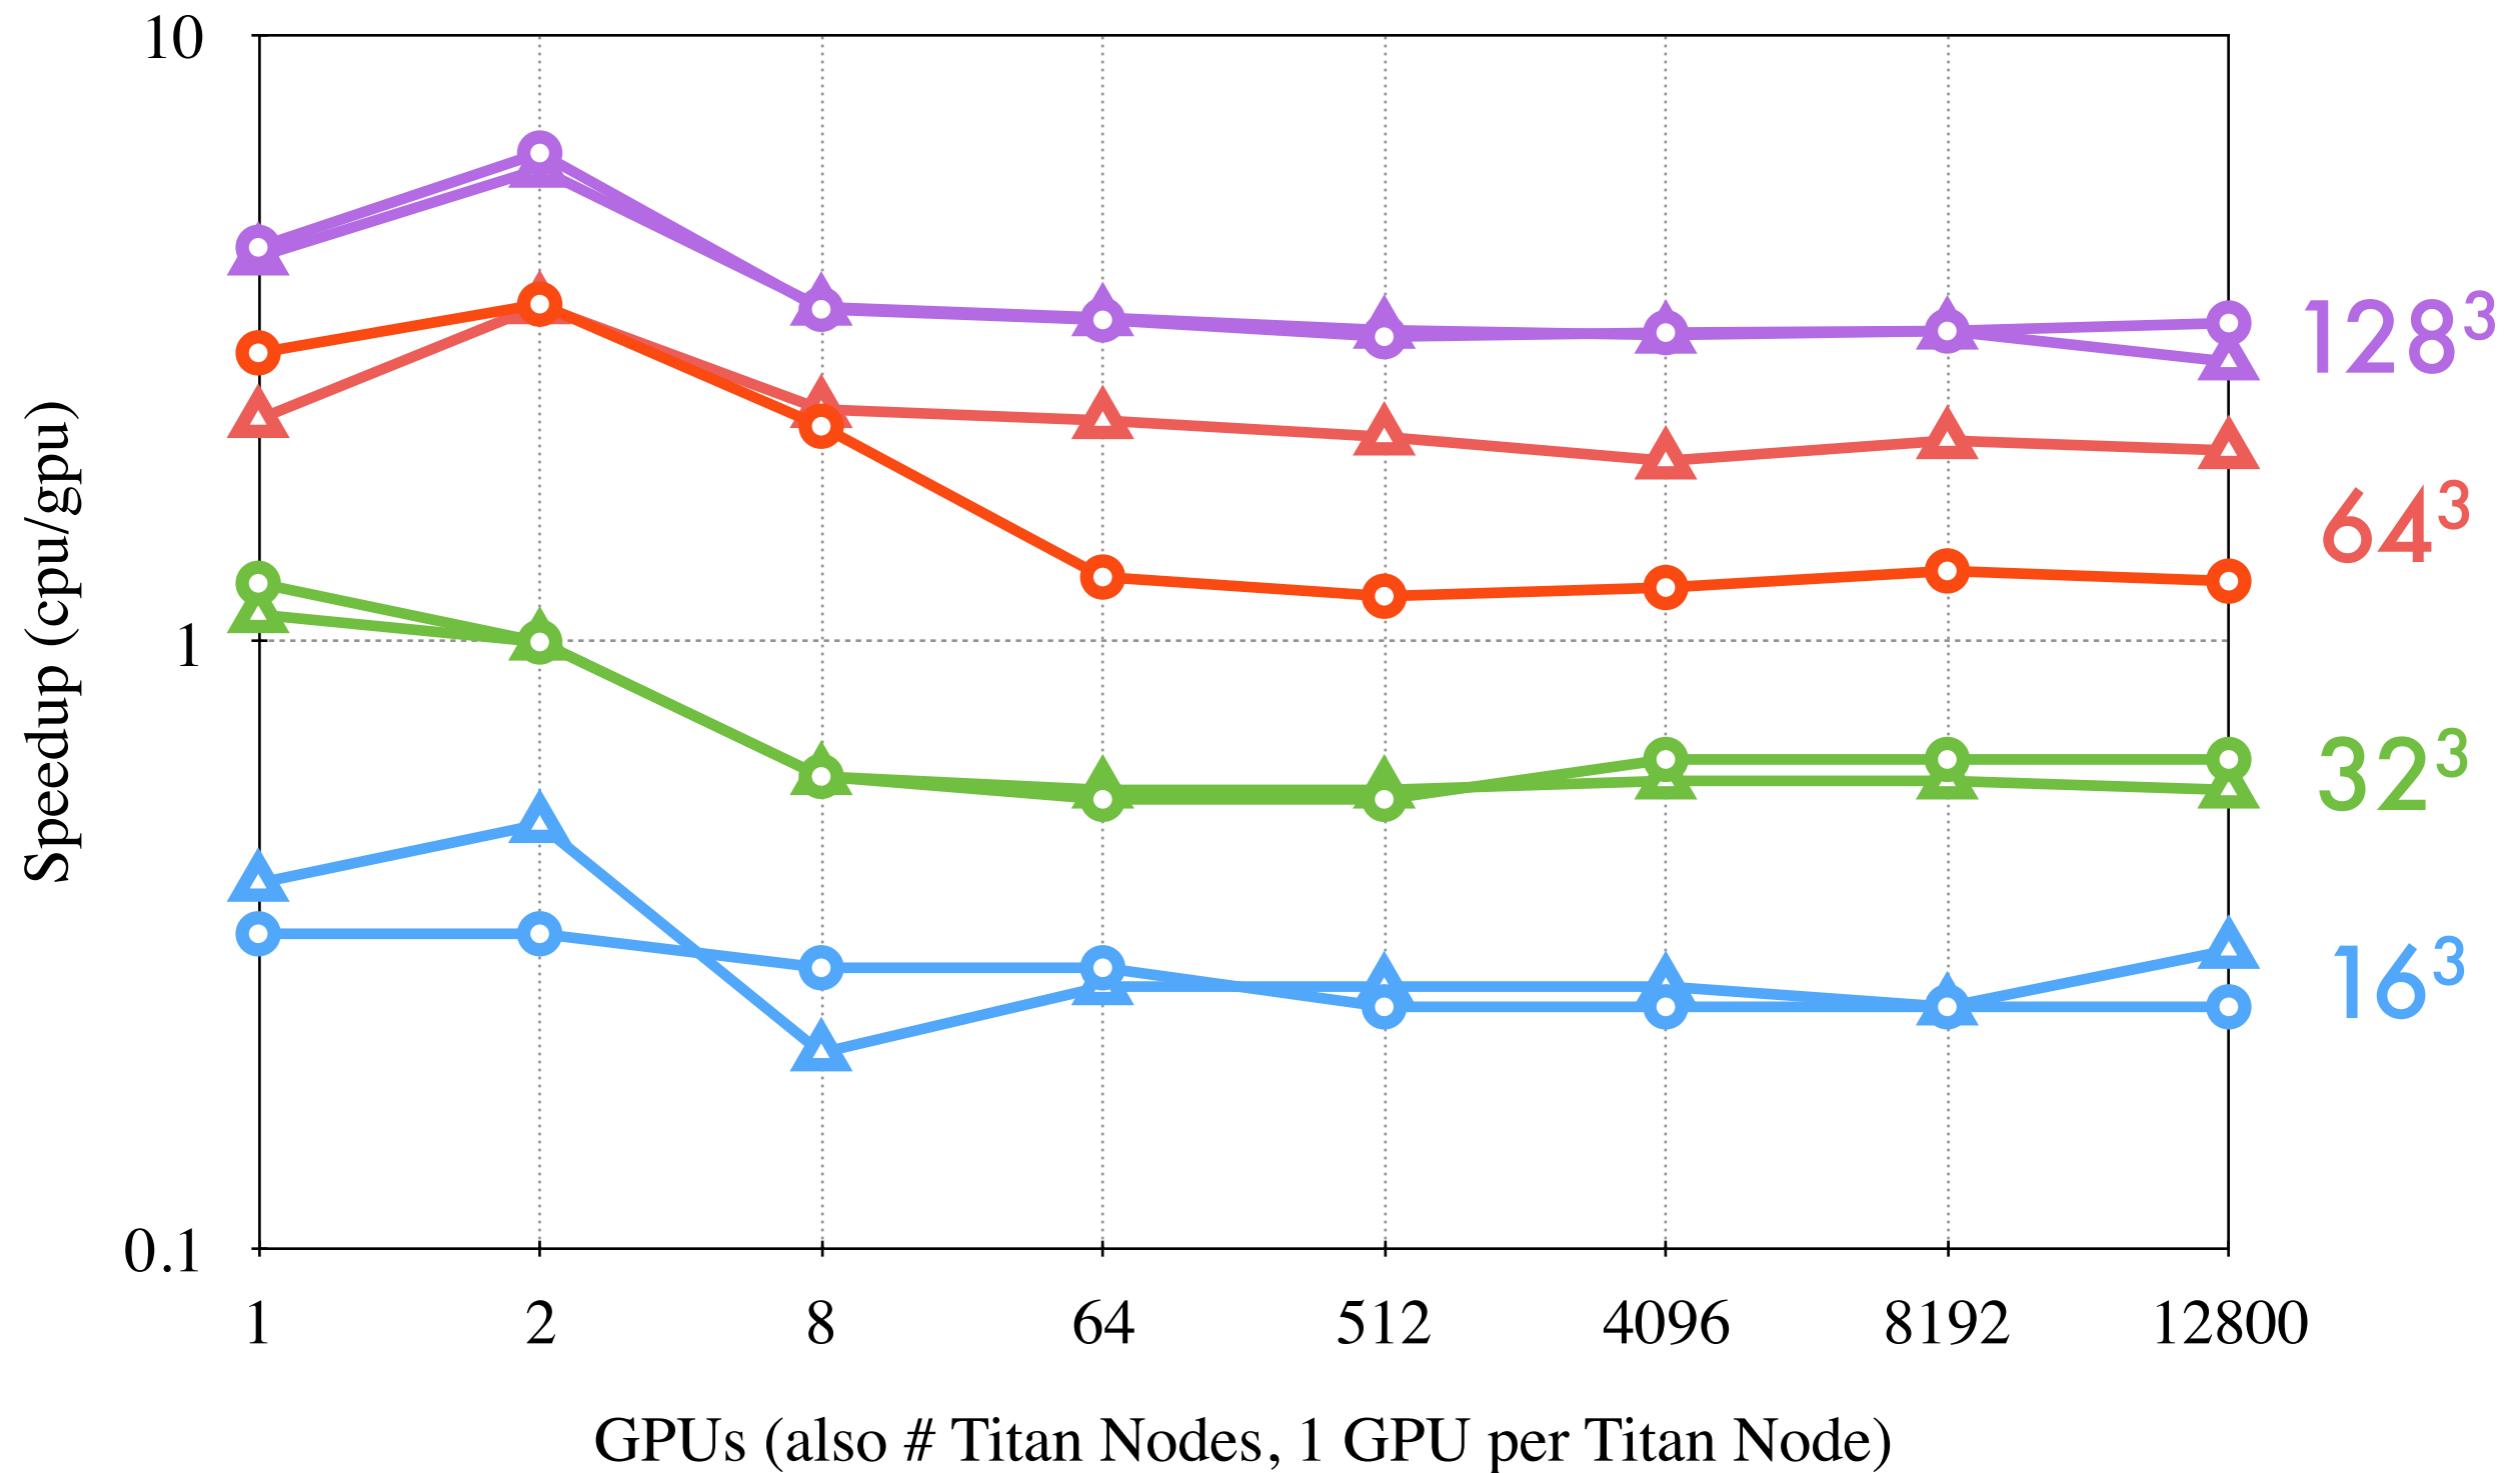
Weak Scaling - Scalar Transport, No Source Terms



○ 10 Equations

△ 20 Equations

Speedup - Scalar Transport, No Source Terms

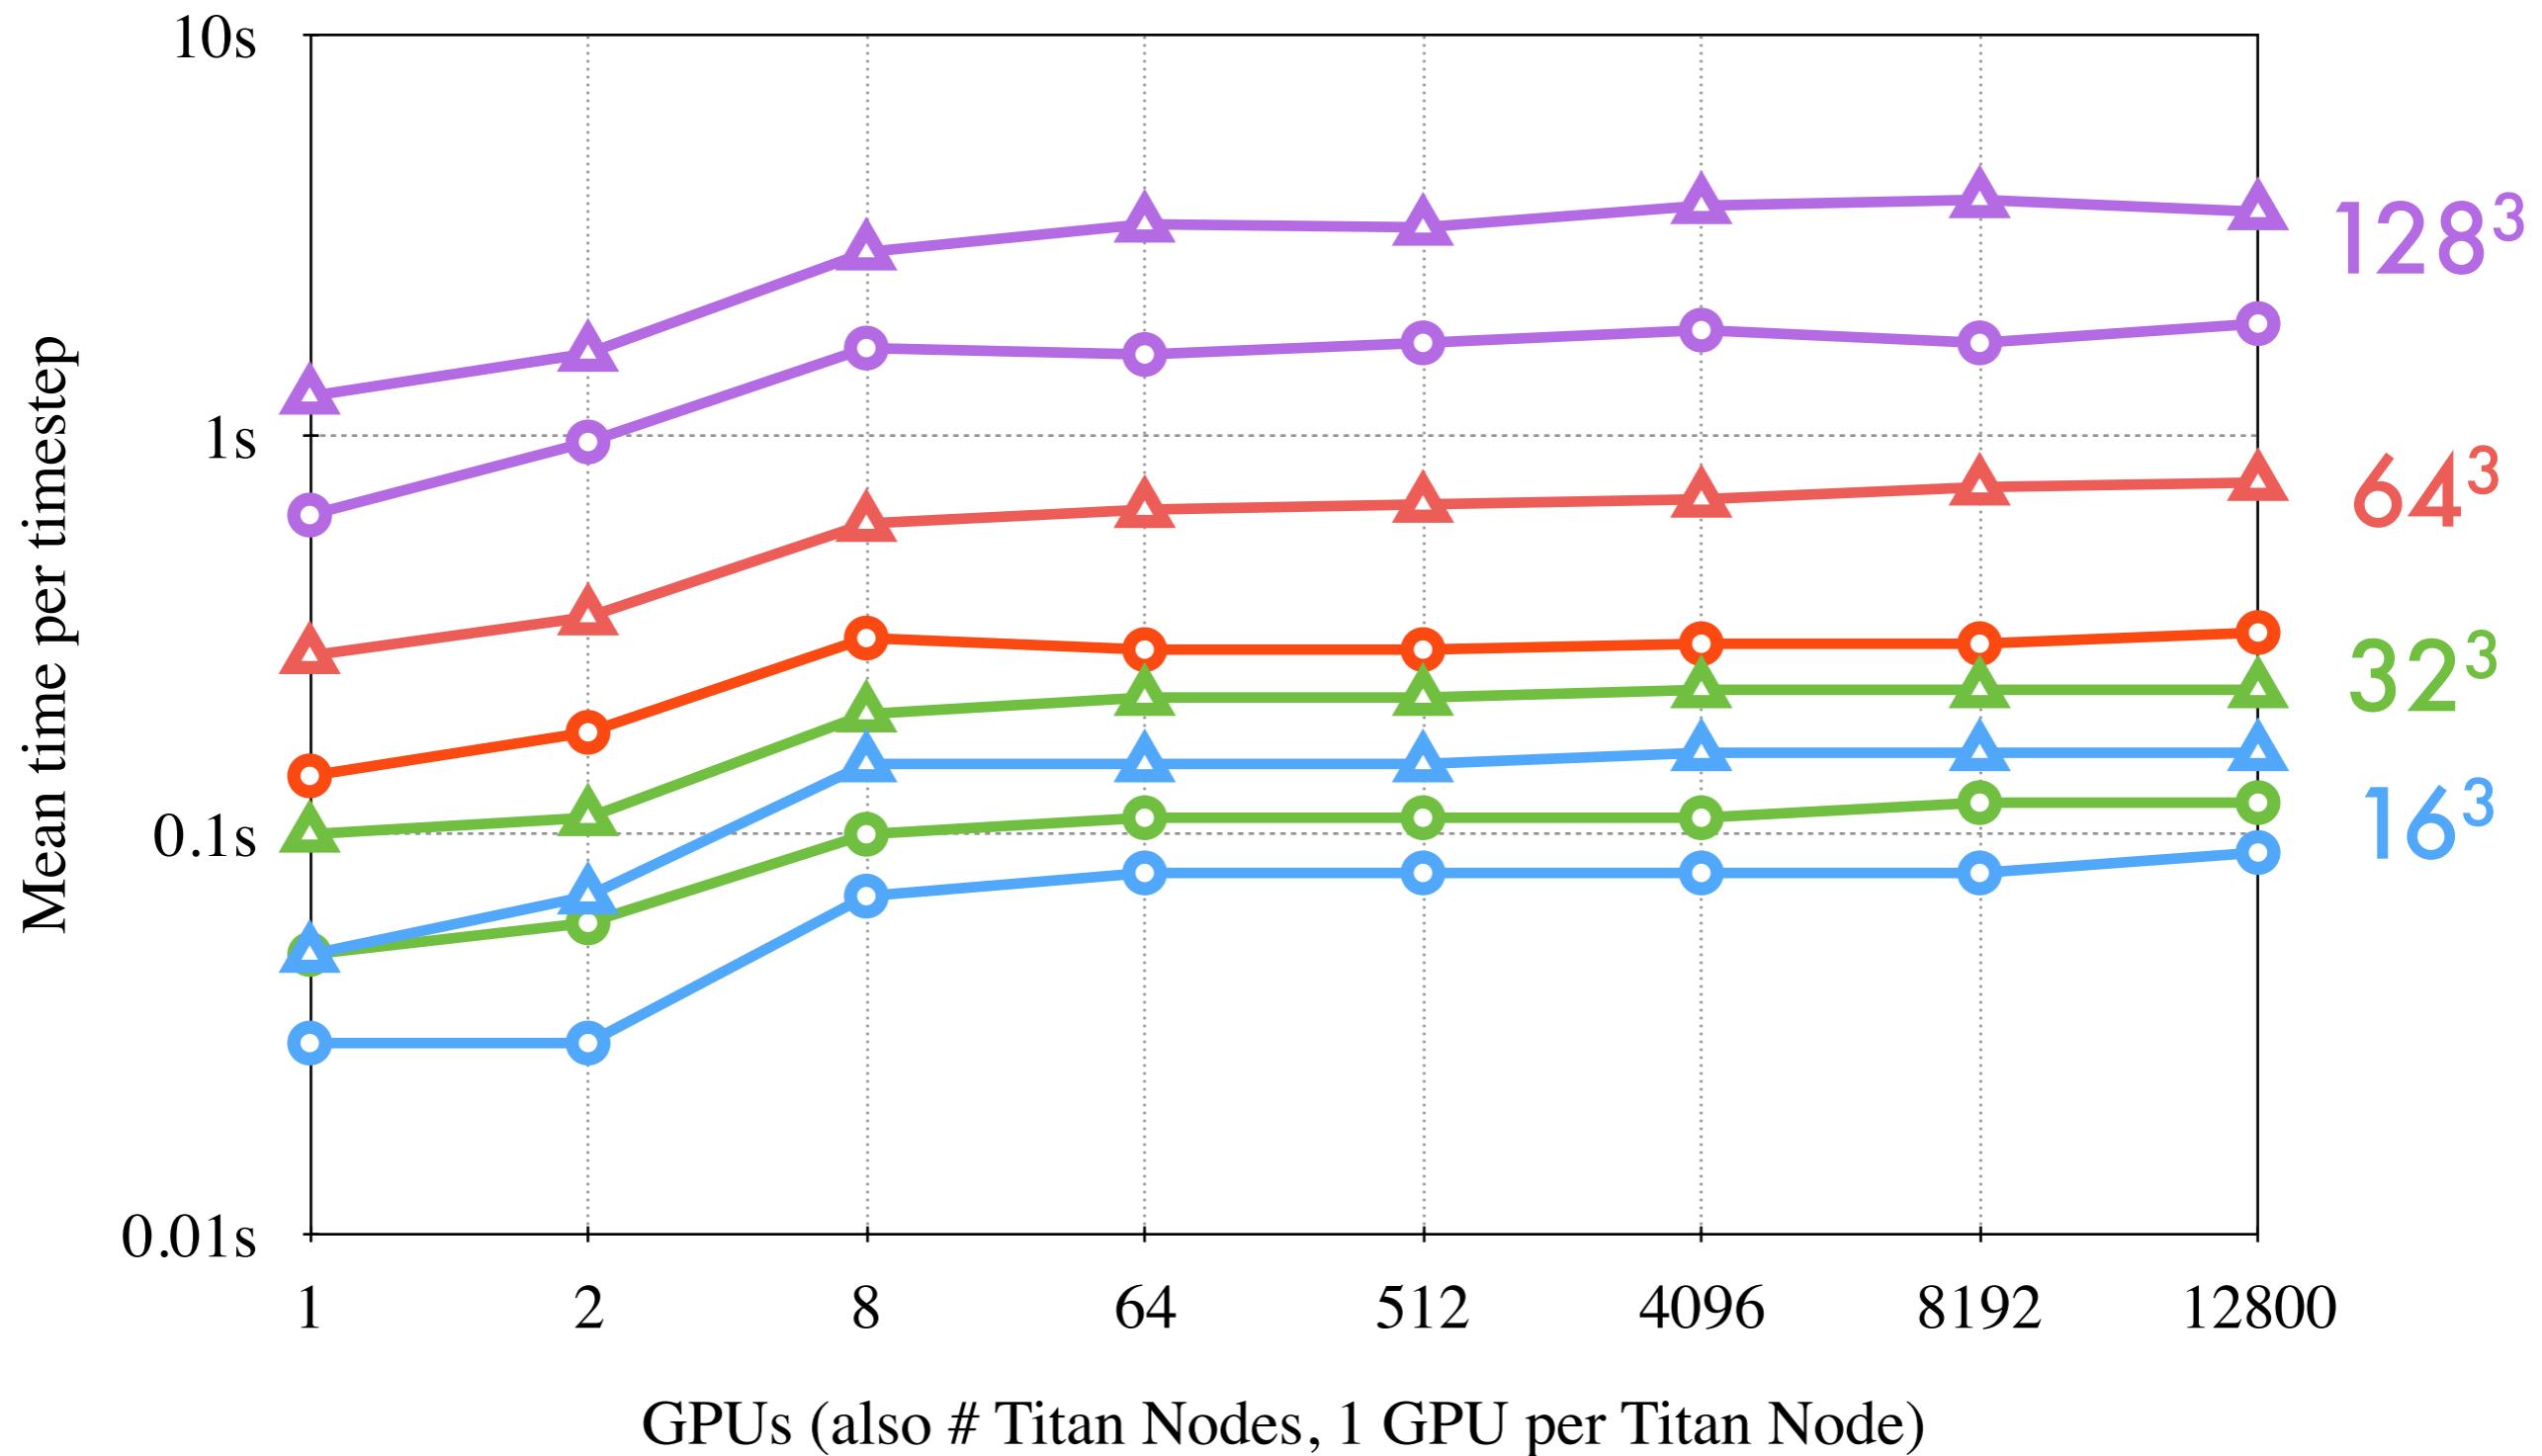


Slightly Complex Physics (Scalar Transport with Uncoupled Source Terms)

○ 10 Equations

△ 20 Equations

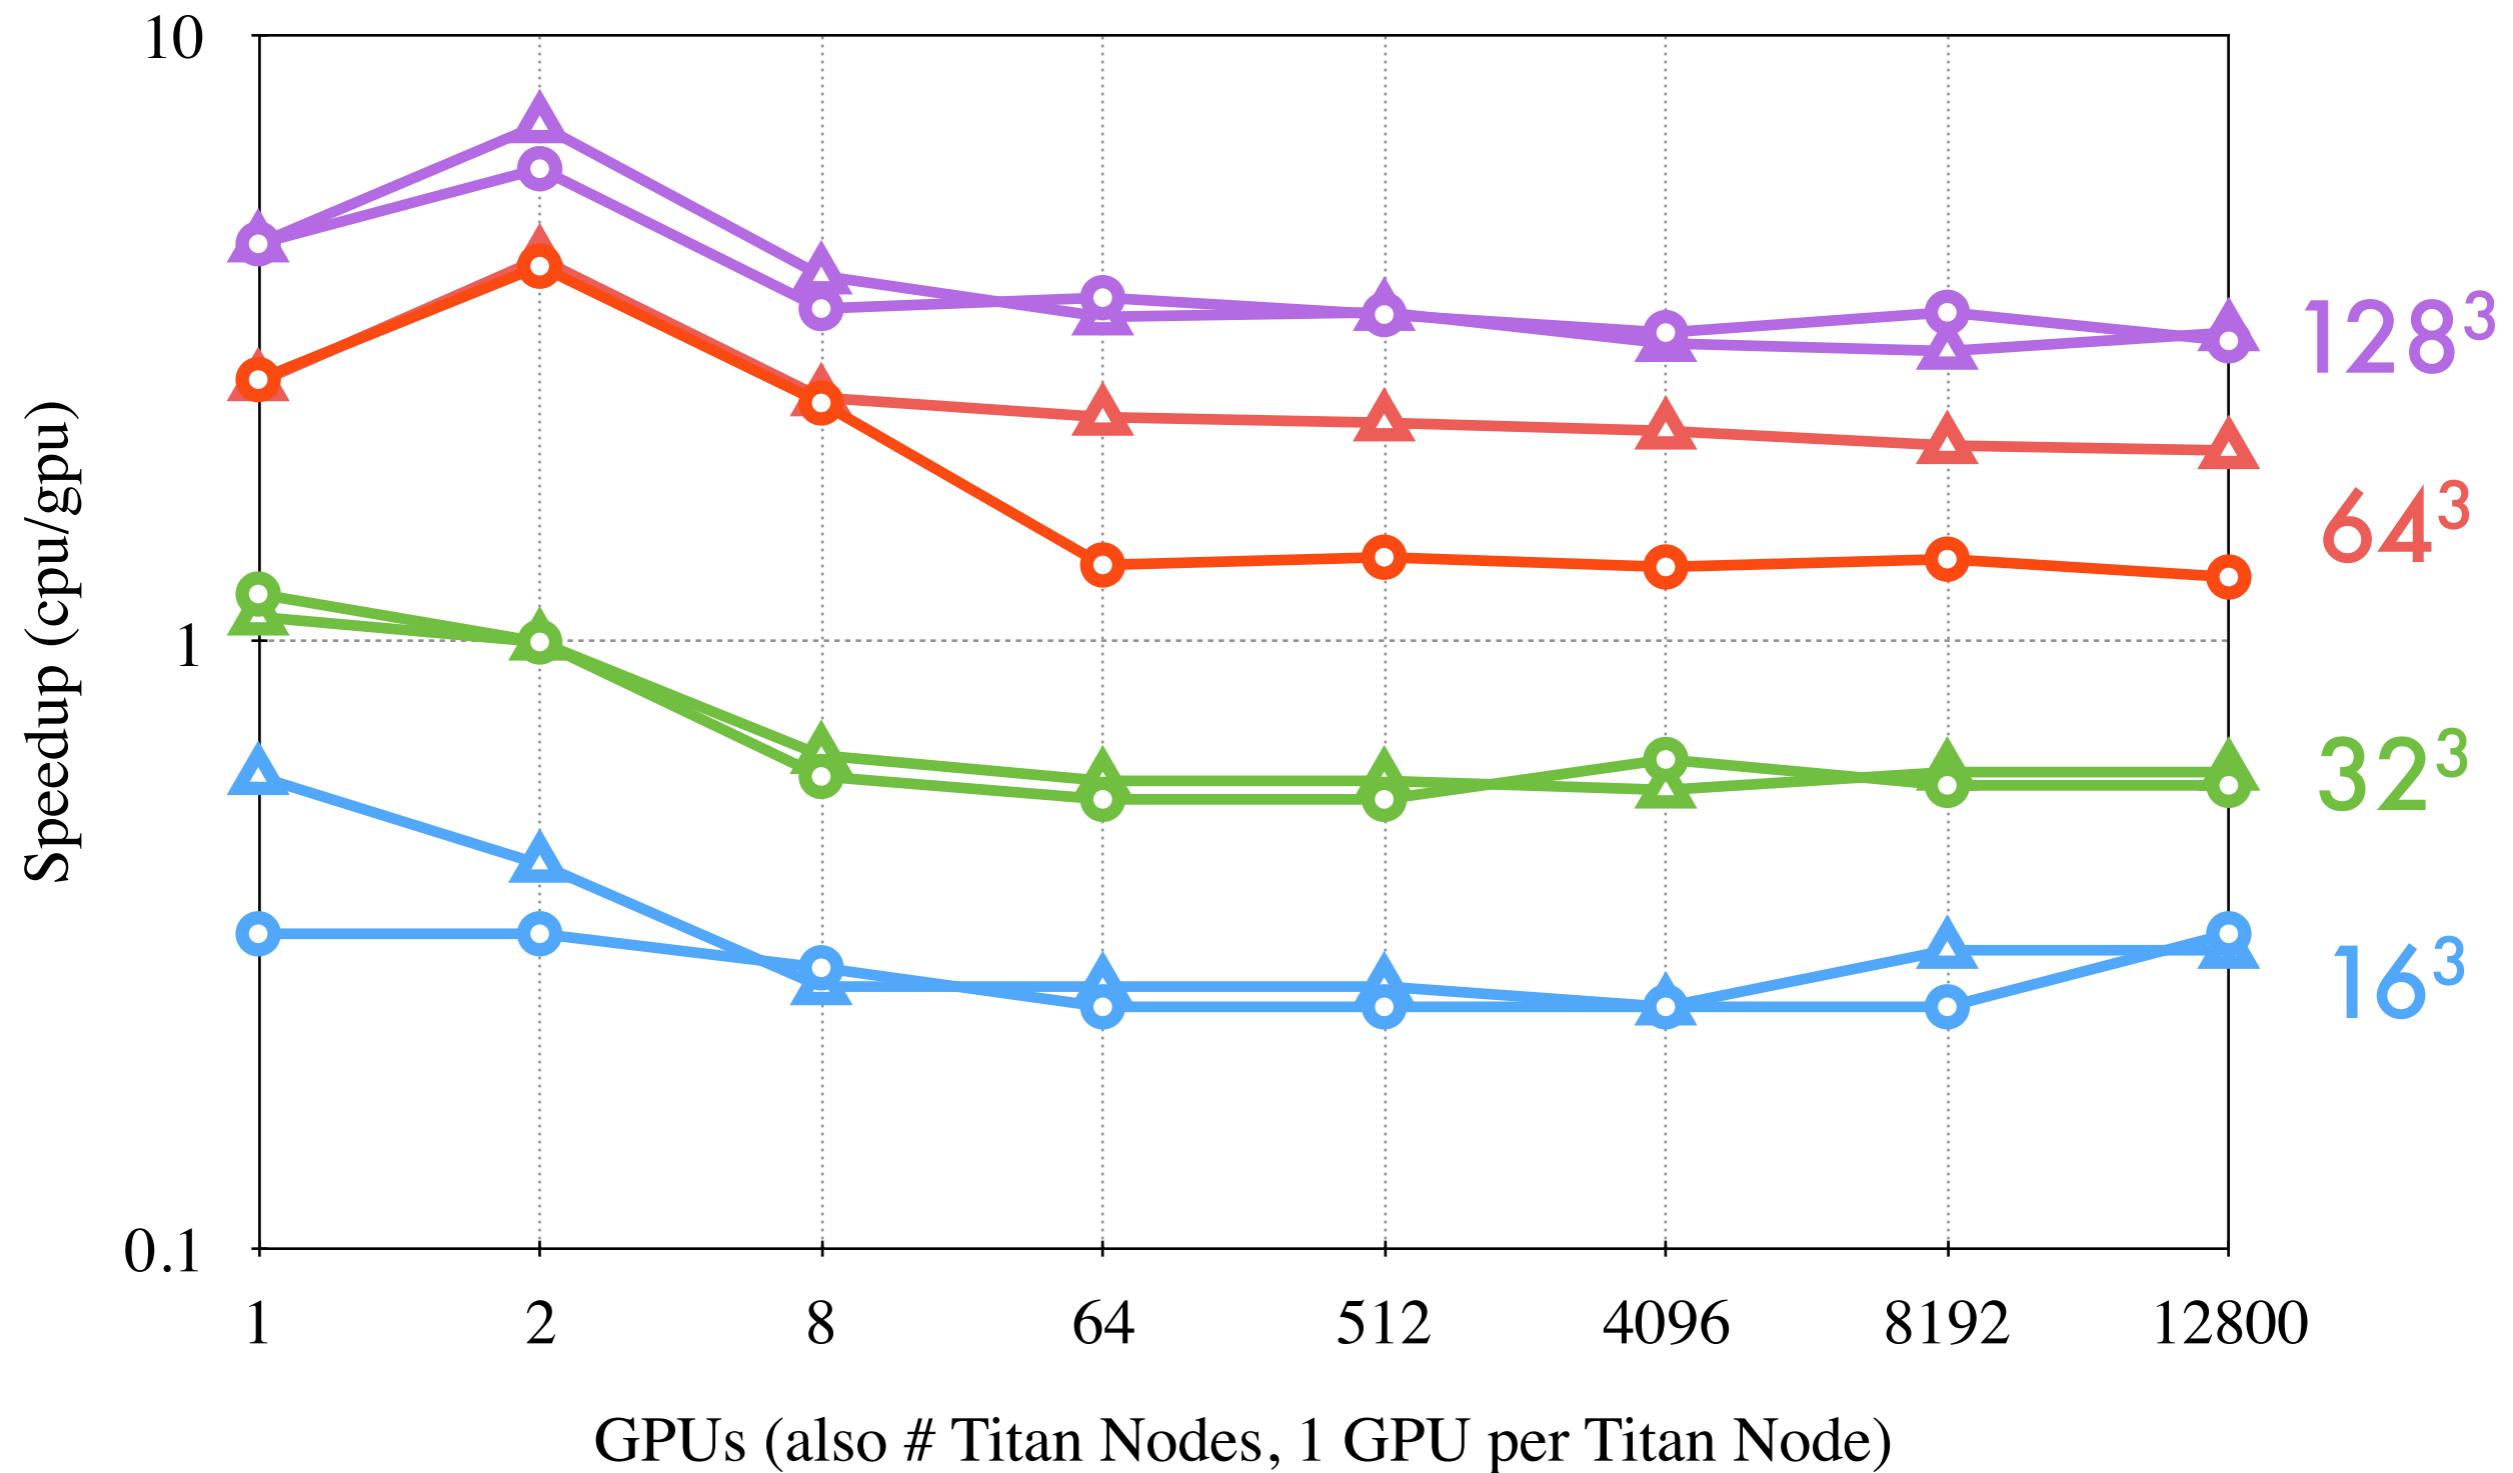
Weak Scaling - Scalar Transport, Uncoupled Source Terms



○ 10 Equations

△ 20 Equations

Speedup - Scalar Transport, No Source Terms

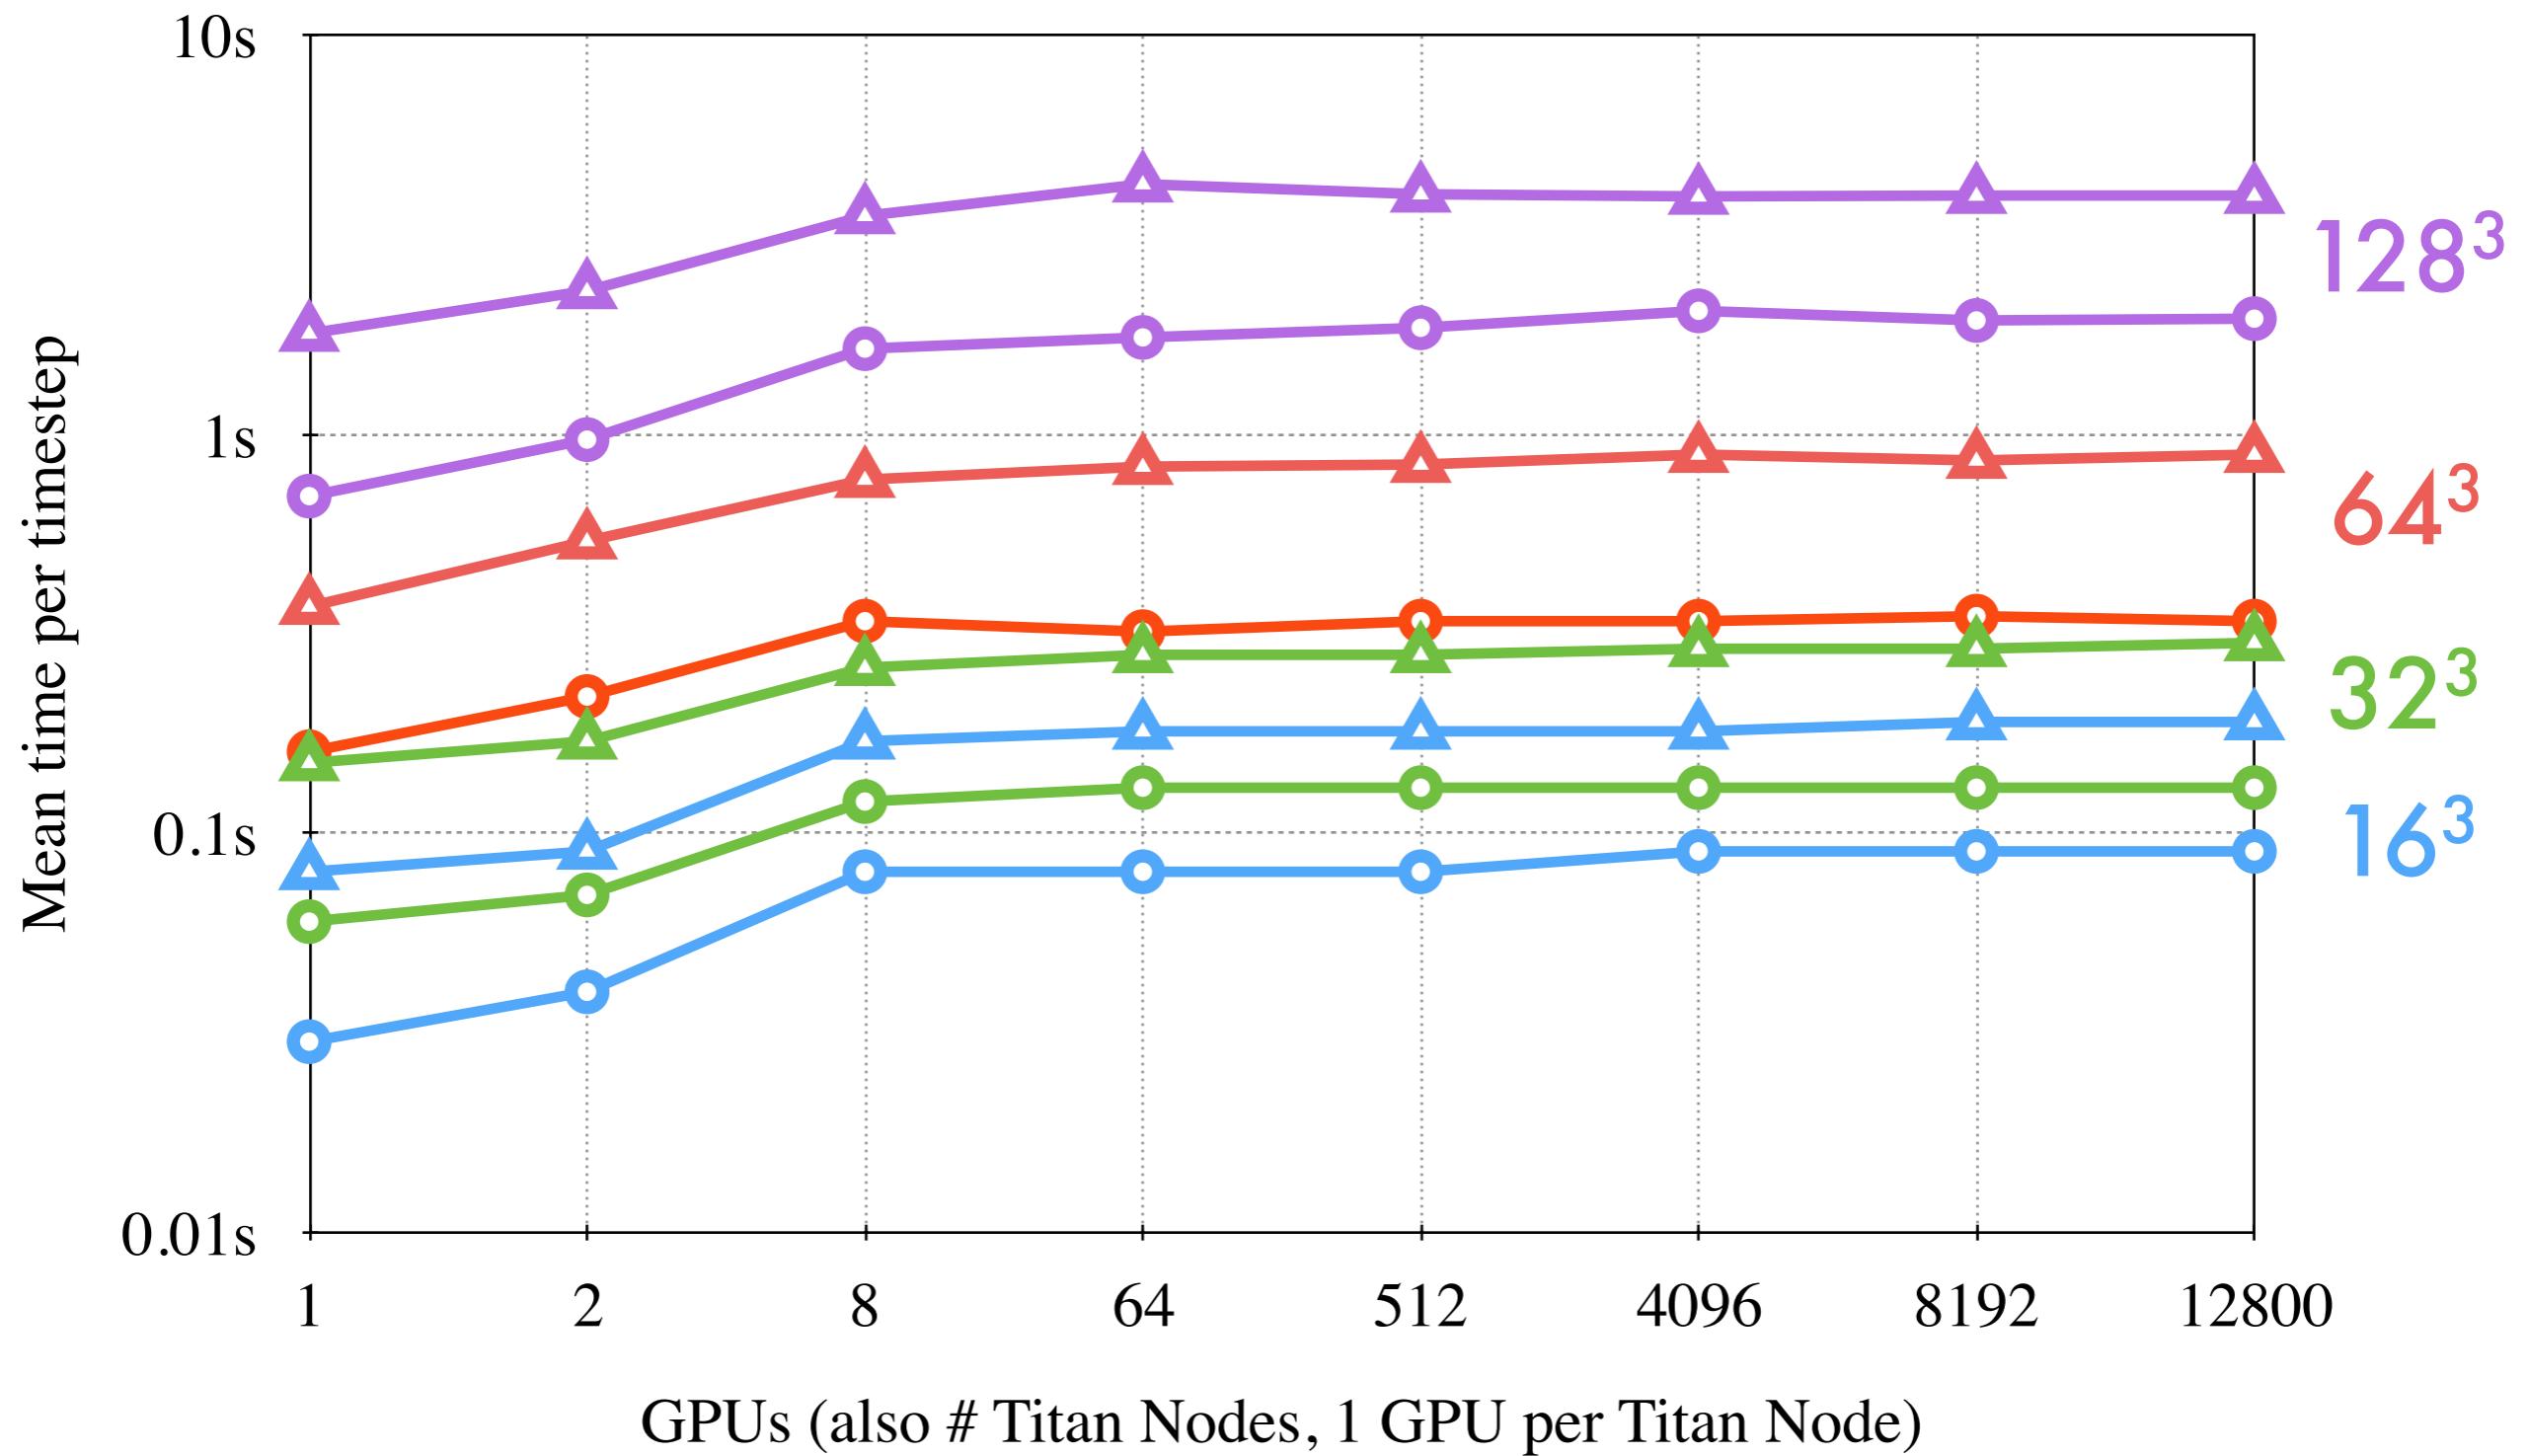


Somewhat Complex Physics (Scalar Transport with Coupled Source Terms)

○ 10 Equations

△ 20 Equations

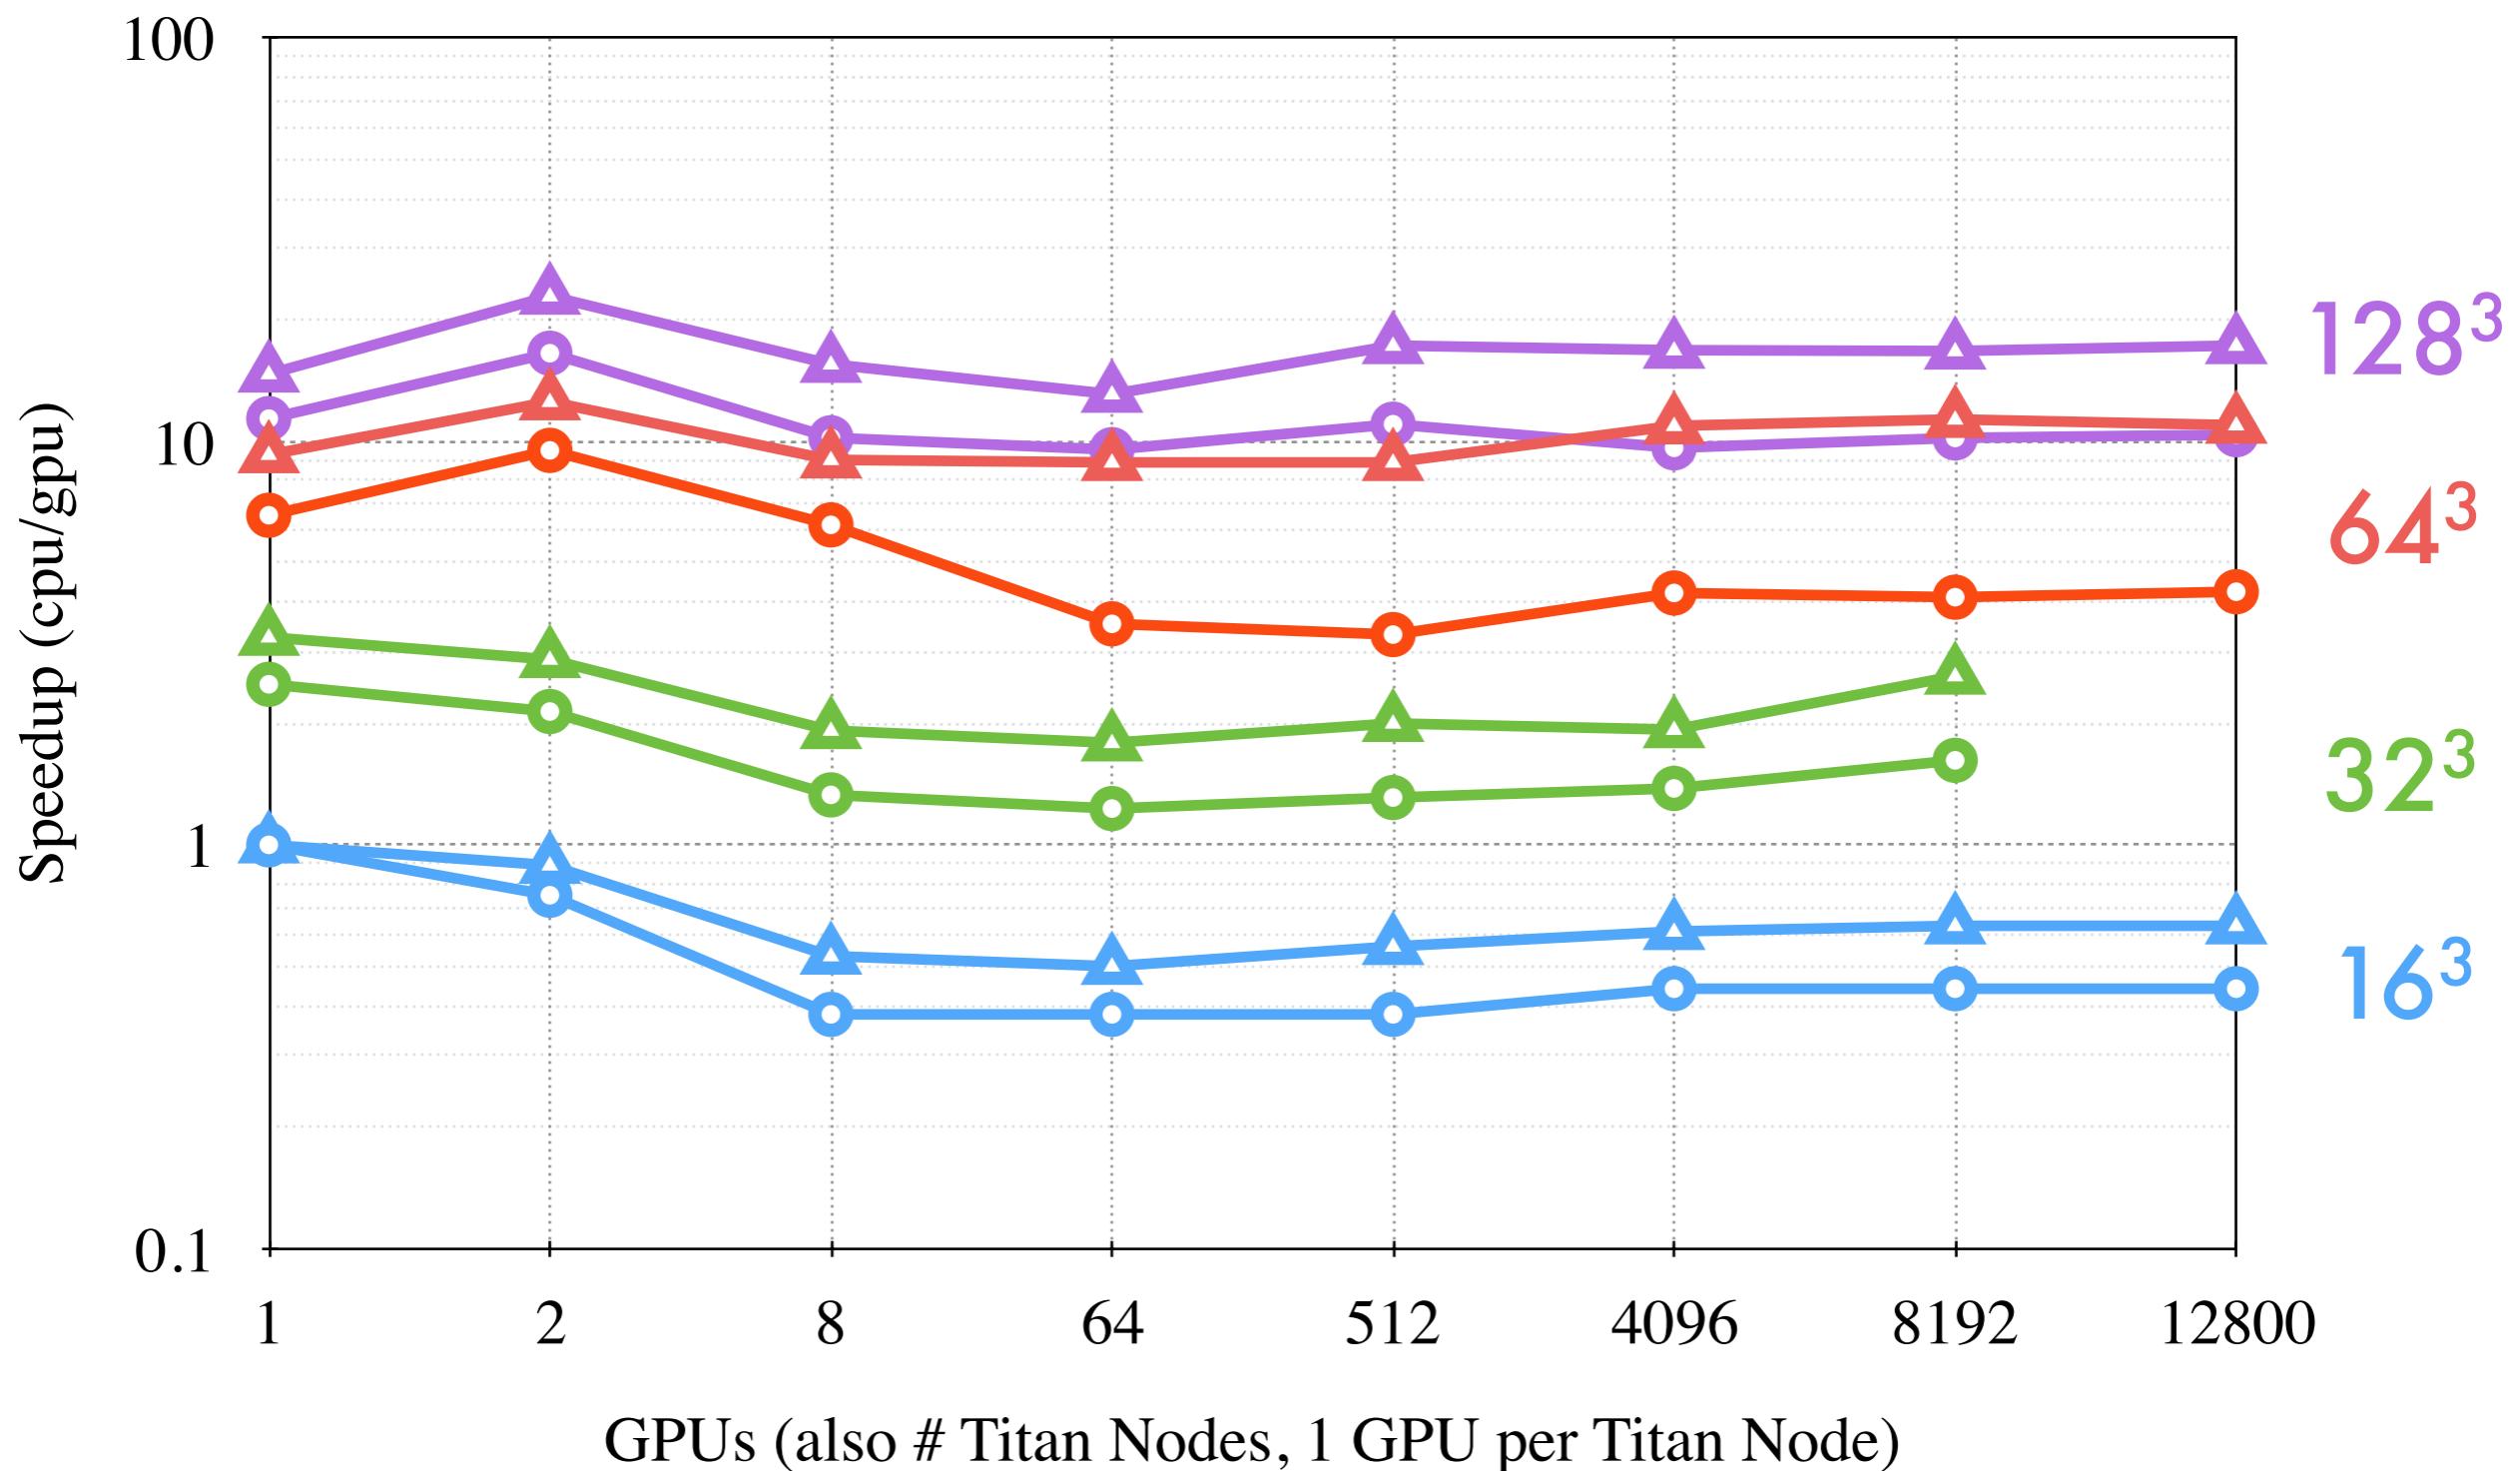
Weak Scaling - Scalar Transport, Coupled Source Terms



○ 10 Equations

△ 20 Equations

Speedup - Scalar Transport, Coupled Source Terms





Q & A

<http://software.crsim.utah.edu>

Tony Saad: tony.saad@utah.edu

James C. Sutherland: james.sutherland@chemeng.utah.edu